

# Introduction: DBMS

©Bipin C. DESAI

Not to be used or distributed in this or any modified form without written permission from the copyright owner.

## Grading:

As per the administrative slides

Only random parts (which could be none) of the assignments could be graded!

*The final letter grade would be based on the traditional conversion scheme; e.g, A's would be assigned to numerical total marks **well above high 80s***

Office hours: TBA in class

For common questions answers would be posted on the announcements in CrsMgr page. Pl. read it before sending an email. Pl. send email in **text**

**DO NOT SEND DUPLICATE emails**

Lab Instructors : TBA announced on CrsMgr

Account for course manager has been created and the ID/PW have been emailed!

For the students who haven't provided an email to the Concordia's SIS, you need to find out your ID/PW!

**Sign in to CrsMgr: change your PW & email address to ENCS**

Course Manager System(CMS):

<https://confsys.encs.concordia.ca/CrsMgr>

Your browser need to accept our self-signed security certificate!

You are already registered in the course manager system.

It would be used for administrating on-line quizzes, managing the grades, submission of assignments and reports, scheduling of demo and peer evaluation for each group marked entity.

Your grade and the feedback could be viewed in CrsMgr.

Please read the announcements for this course regularly in CrsMgr.

Answers to common questions as well additional instructions for the course would be posted on the flash page.

Look at it before sending out emails.

No emails will be answered if the instructions are already posted.

Pl. send email in text – **DOES NOT MEAN TEXT MESSAGES**

**Sign in to CrsMgr and change your email adr. to the ENCS email.**

Form a team of 1-4 in the first class or before the deadline:

Choose a leader: the leader signs into CrsMgr and joins a new group

- update his/her email to **ENCS** computer account
- ~~insert a password for group DB (unless generated automatically):~~

Each member of the team then joins this group.

For each member of the team:

- update his/her email to **ENCS** computer account

This could be done by using and uploading a text file!

Those who do not join a group will be put randomly in a group to create groups with up to 4 members

Note: Assignments and Projects (warm up and main) would be done by this group with a single submission per group to be uploaded by the group(team) leader.

Upload the assignment/project to the course manager system

<https://confsys.encs.concordia.ca/CrsMgr>

# Software

On the ENCS system:

MySQL (Oracle is no longer supported by AITS!)

PHP

HTML

CGI (security) – defunct: each group have their own system

**On you own( for WinX or Linux) NO HELP from us:**

Download and install MySQL

<http://dev.mysql.com/>

Install PHP

<http://www.php.net/>

Install Web Server

<http://www.apache.org/>

Look into LAMP or its  
port to Windows

**Remember the demo would be on the ENCS system so you need to port it.**

If you use other database engines:

**NO HELP from us**

Only MySQLis supported by ENCS's AITS

You are free to use any database engine.

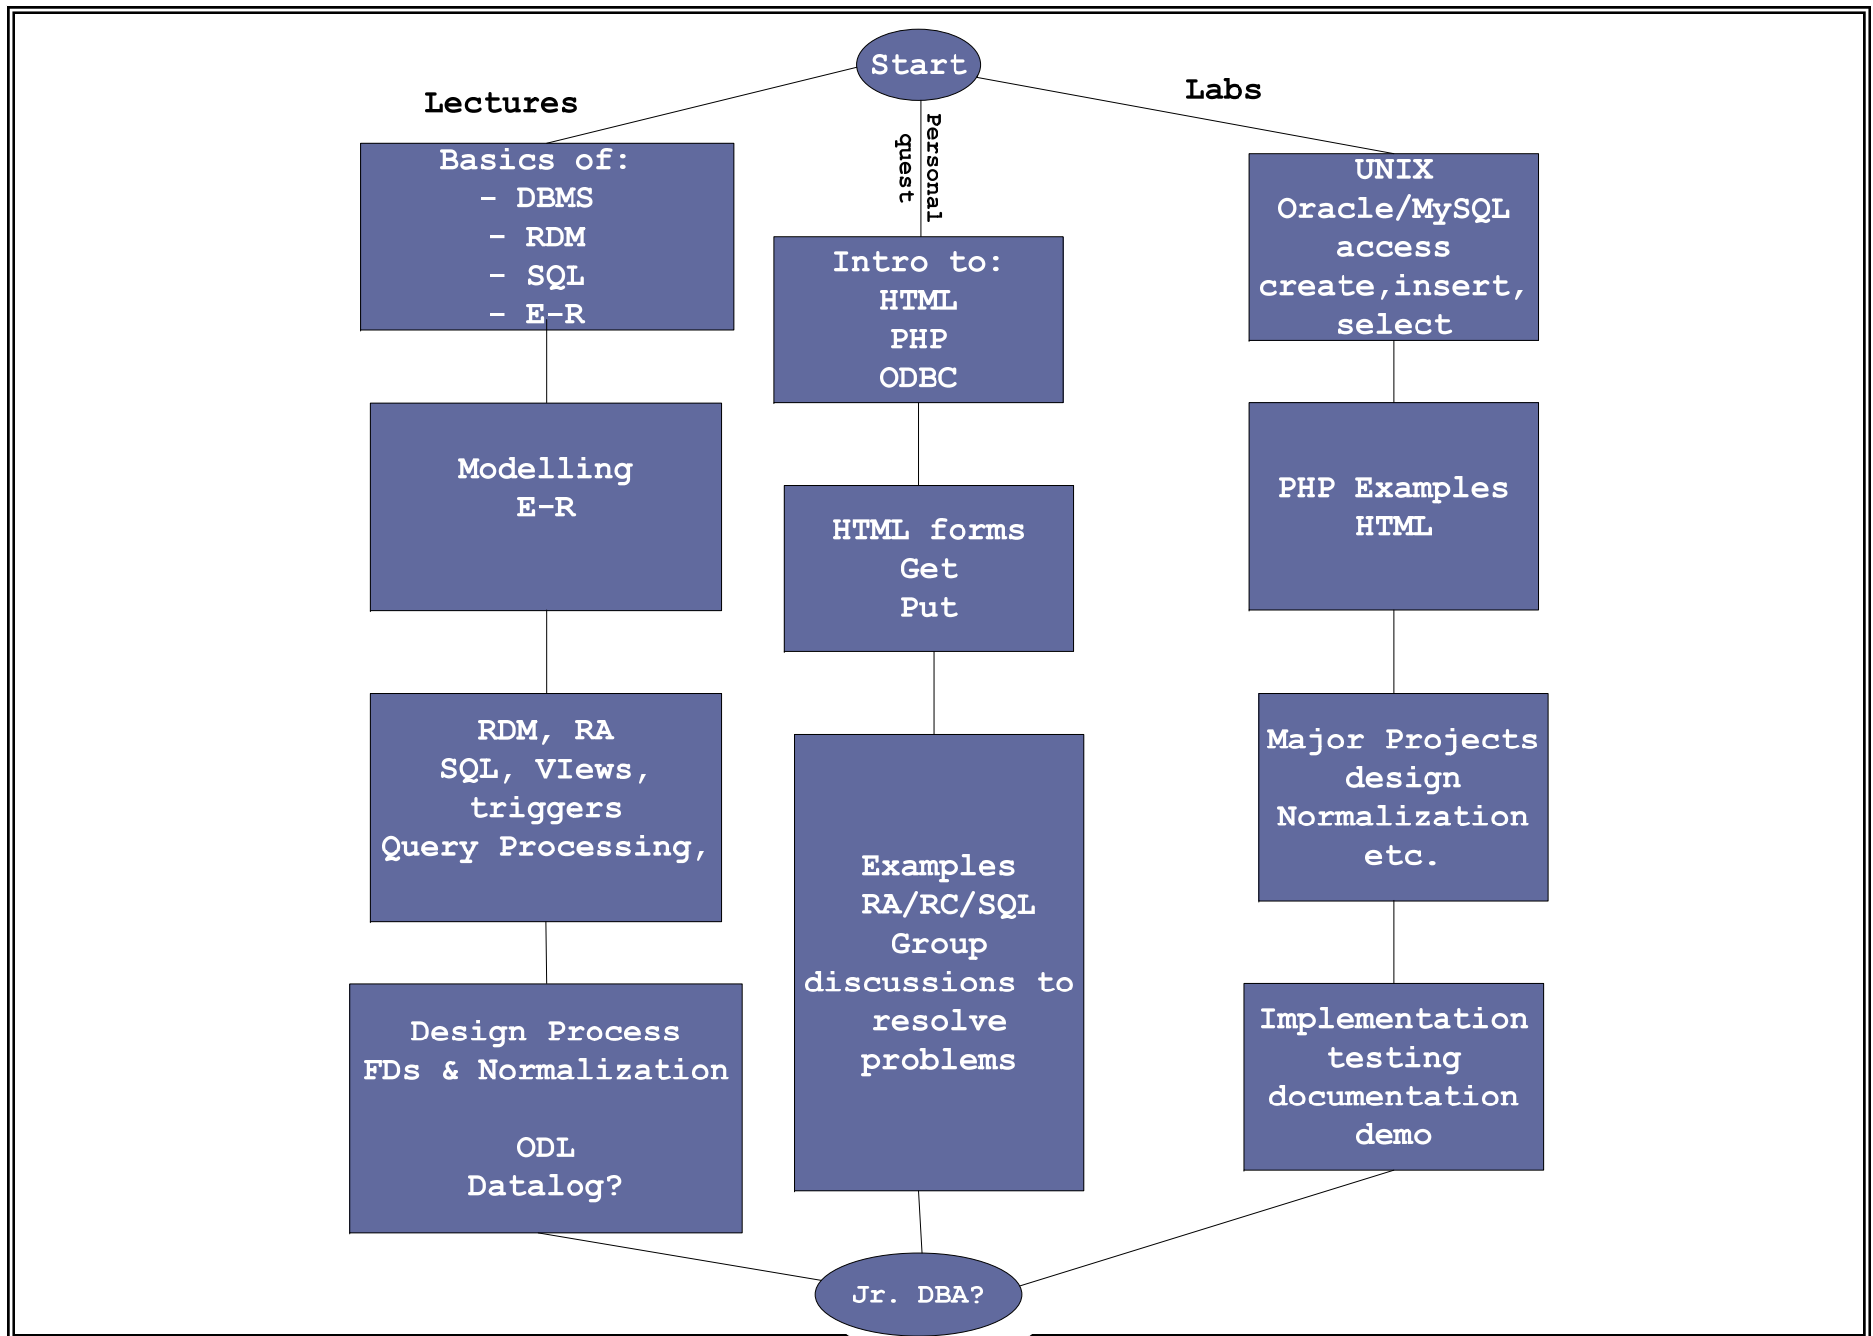
However, you are on your own to download and install the database on your own system;

You also need to make arrangements with your teammates so that the work is coordinated.

You need to install PHP and Apache servers as well.

The database system applications you develop for the projects must be compatible with Linux (optionally WinX) and ported to one of our system for the demo.

**Remember the demo would be on the ENCS system so you need to port your projects!**



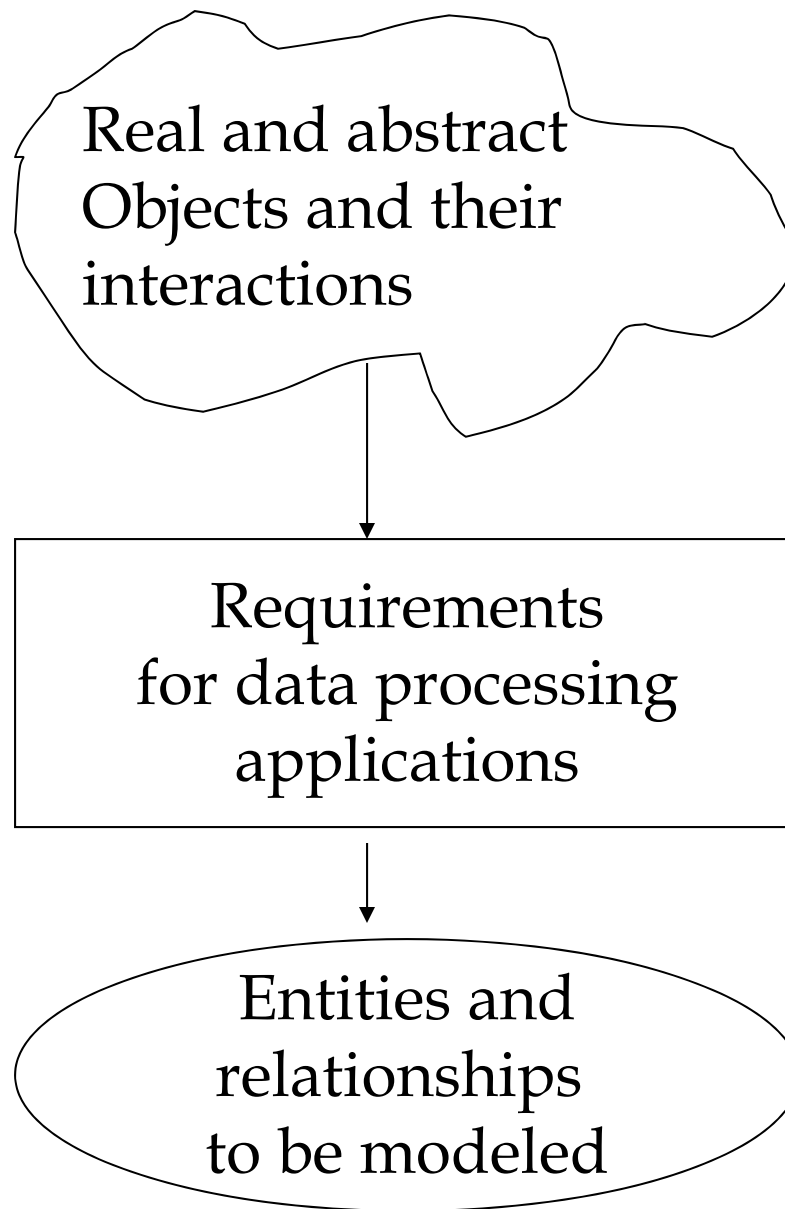


Modeling techniques ( E-R, ODL)

Basic relational model

Design of database applications

Database programming (MySQL, SQL, PHP,  
HTML)



# DBMS! What is it?

Database is an integrated data collection  
(Logically consistent and persistent)

It is derived from the model of a set of applications for a real world enterprise.

DBMS is a software package designed to make managing almost any database.

DBMS offers: data independence, efficiency, integrity, security, concurrency, recovery

# Why Database?

Information Age: 30-40% of world trade and growing

Web(Unstructured data) and .com

Digital Library

Human Genome Project

Day to day operation of Mama/Papa Store

List of titles, artist, and download site of shared files.

Information about employees, departments, projects,  
etc. in an organization

Information about students, courses, enrollments,  
professors, etc. in an educational institute

Information about books, videos, albums, members, etc.  
in a library

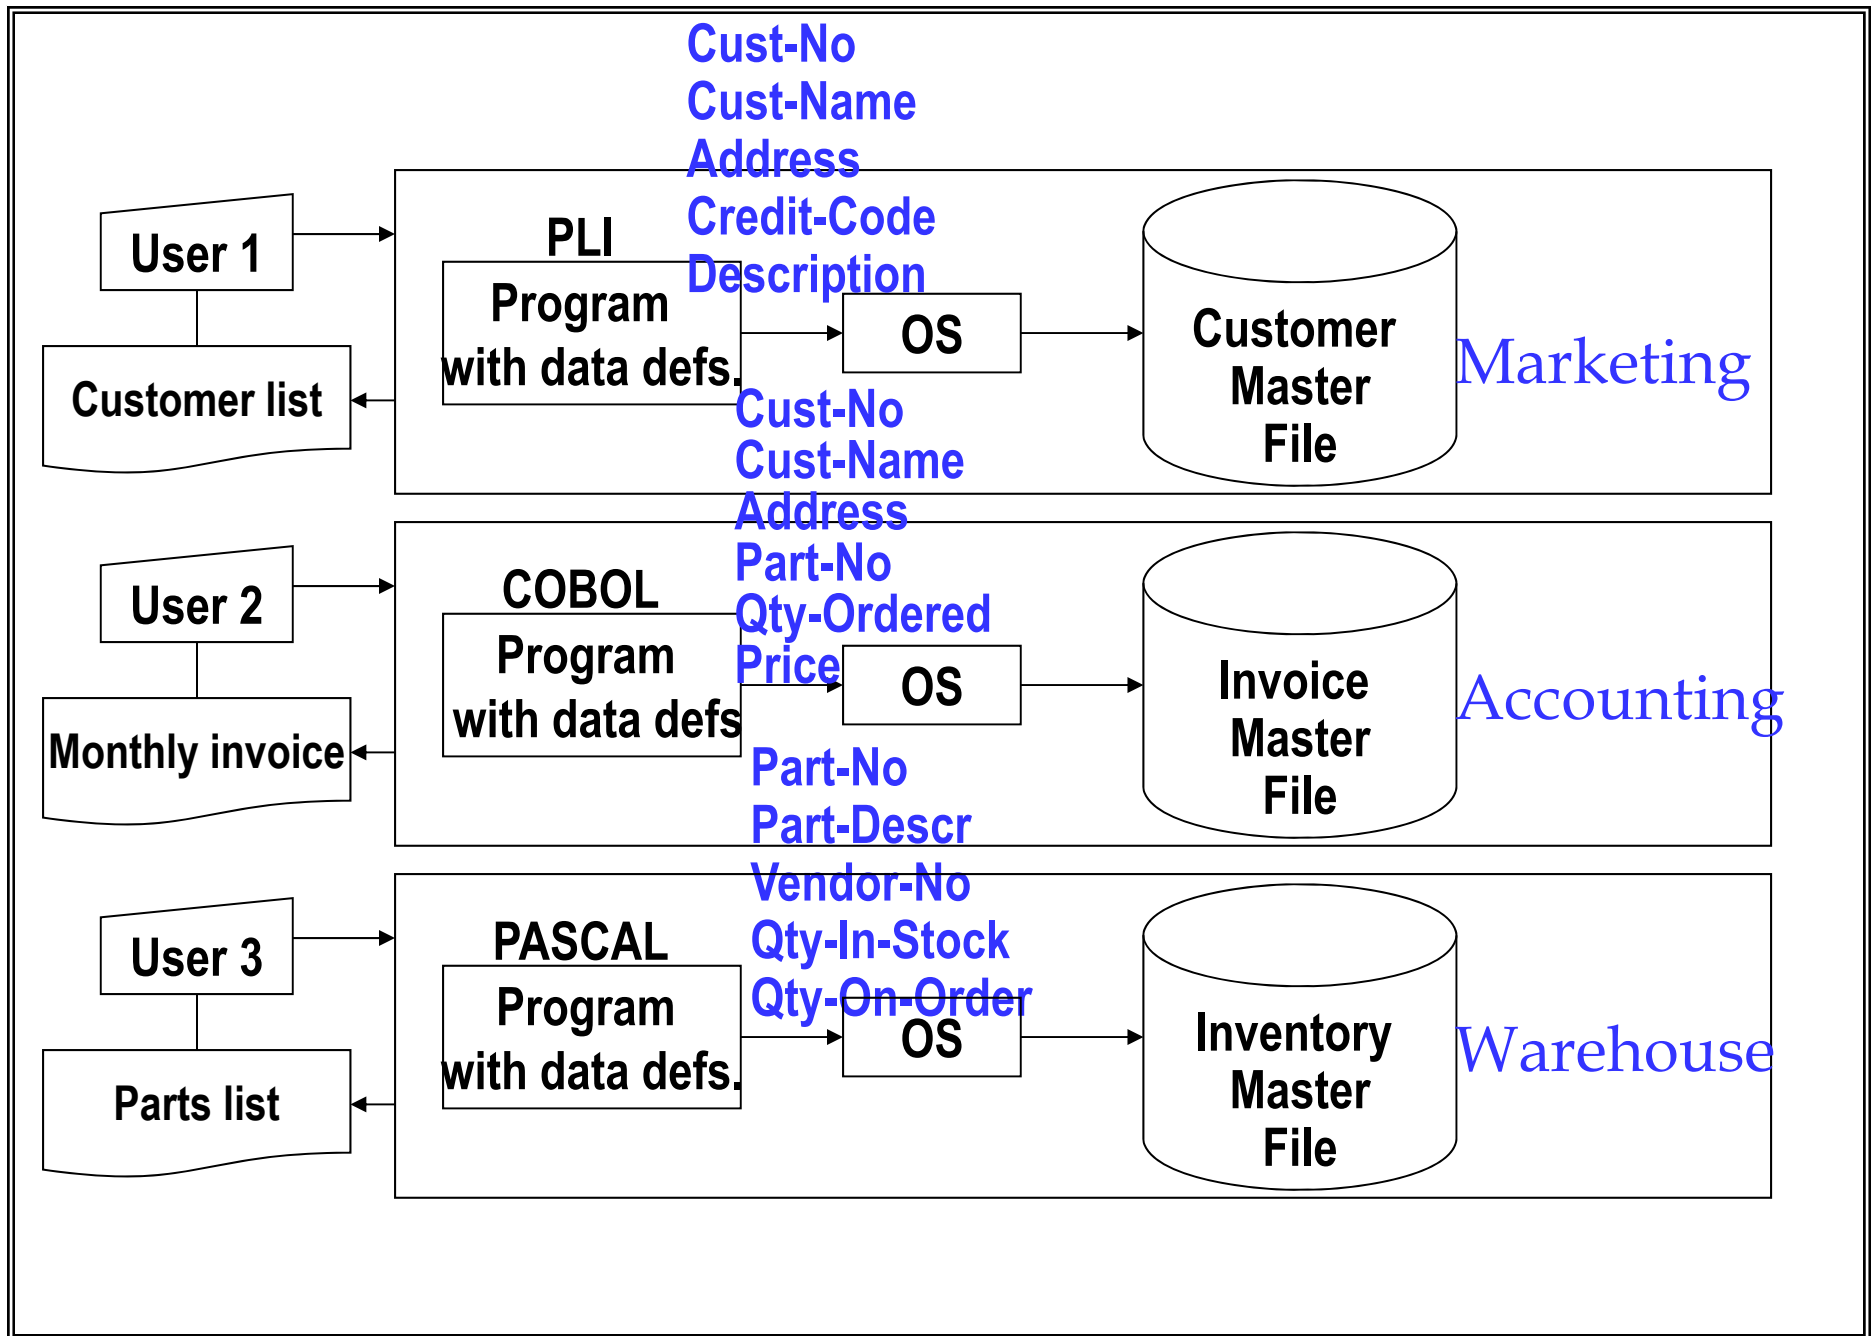
DBMS is a complex set of software packages:

- create new databases, store and manage data
- provide application development and support environment

**Application Support:** Gives developers tools to build applications for using the data. Allows easy method for users to query and modify the data

**Persistent storage:** Support the storage of data

**Transaction management:** Controls concurrent access to data from many users (supports the Atomicity Concurrency Integrity Durability properties ACID.)



# Pros & Cons of file based system

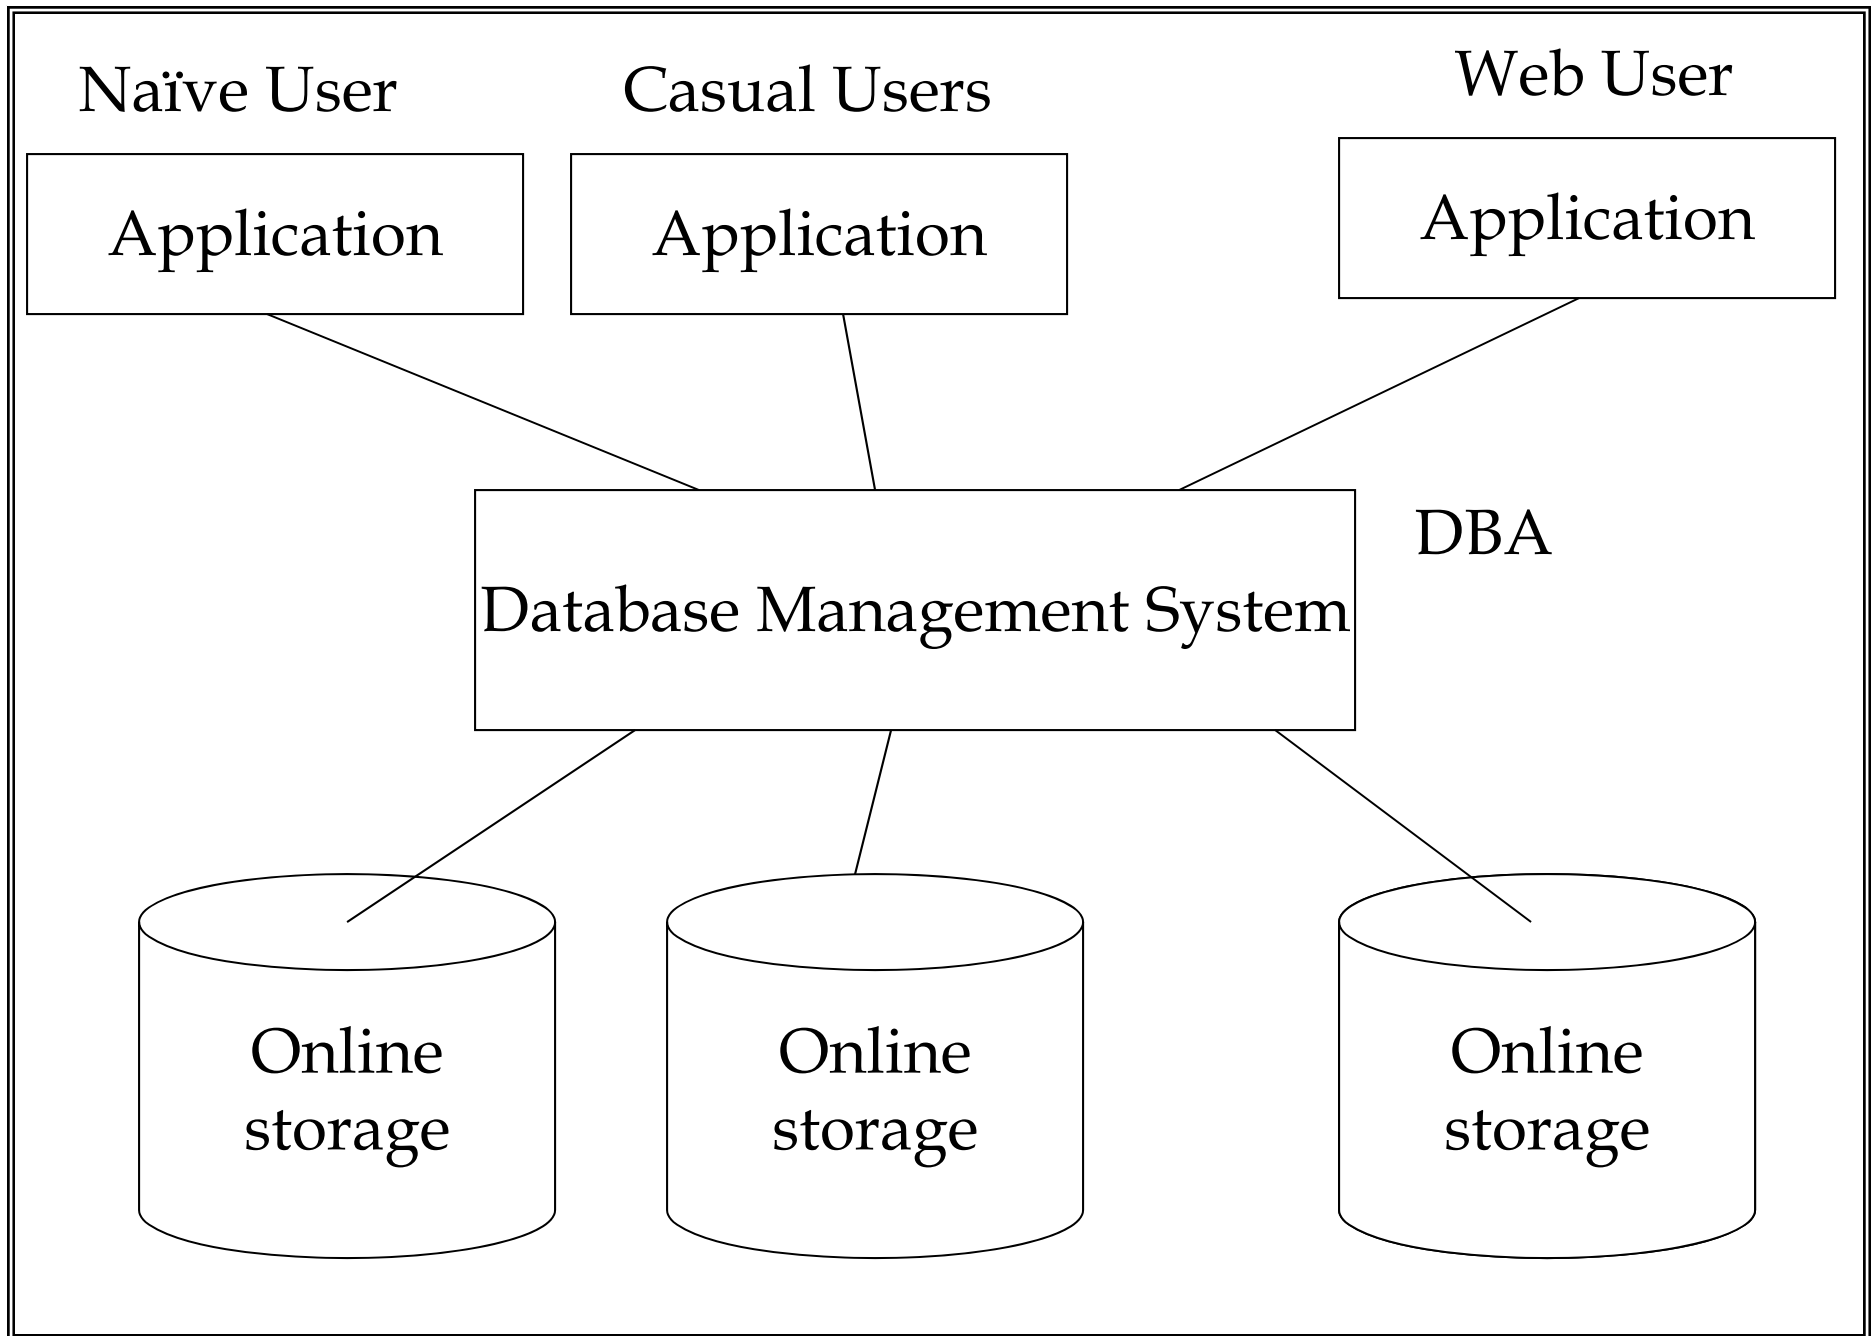
Sharing not possible data definition is “locked” in application programs which “owns” the file and the data in it

Redundancy of data: Same data is duplicated perhaps in slightly different format over various files

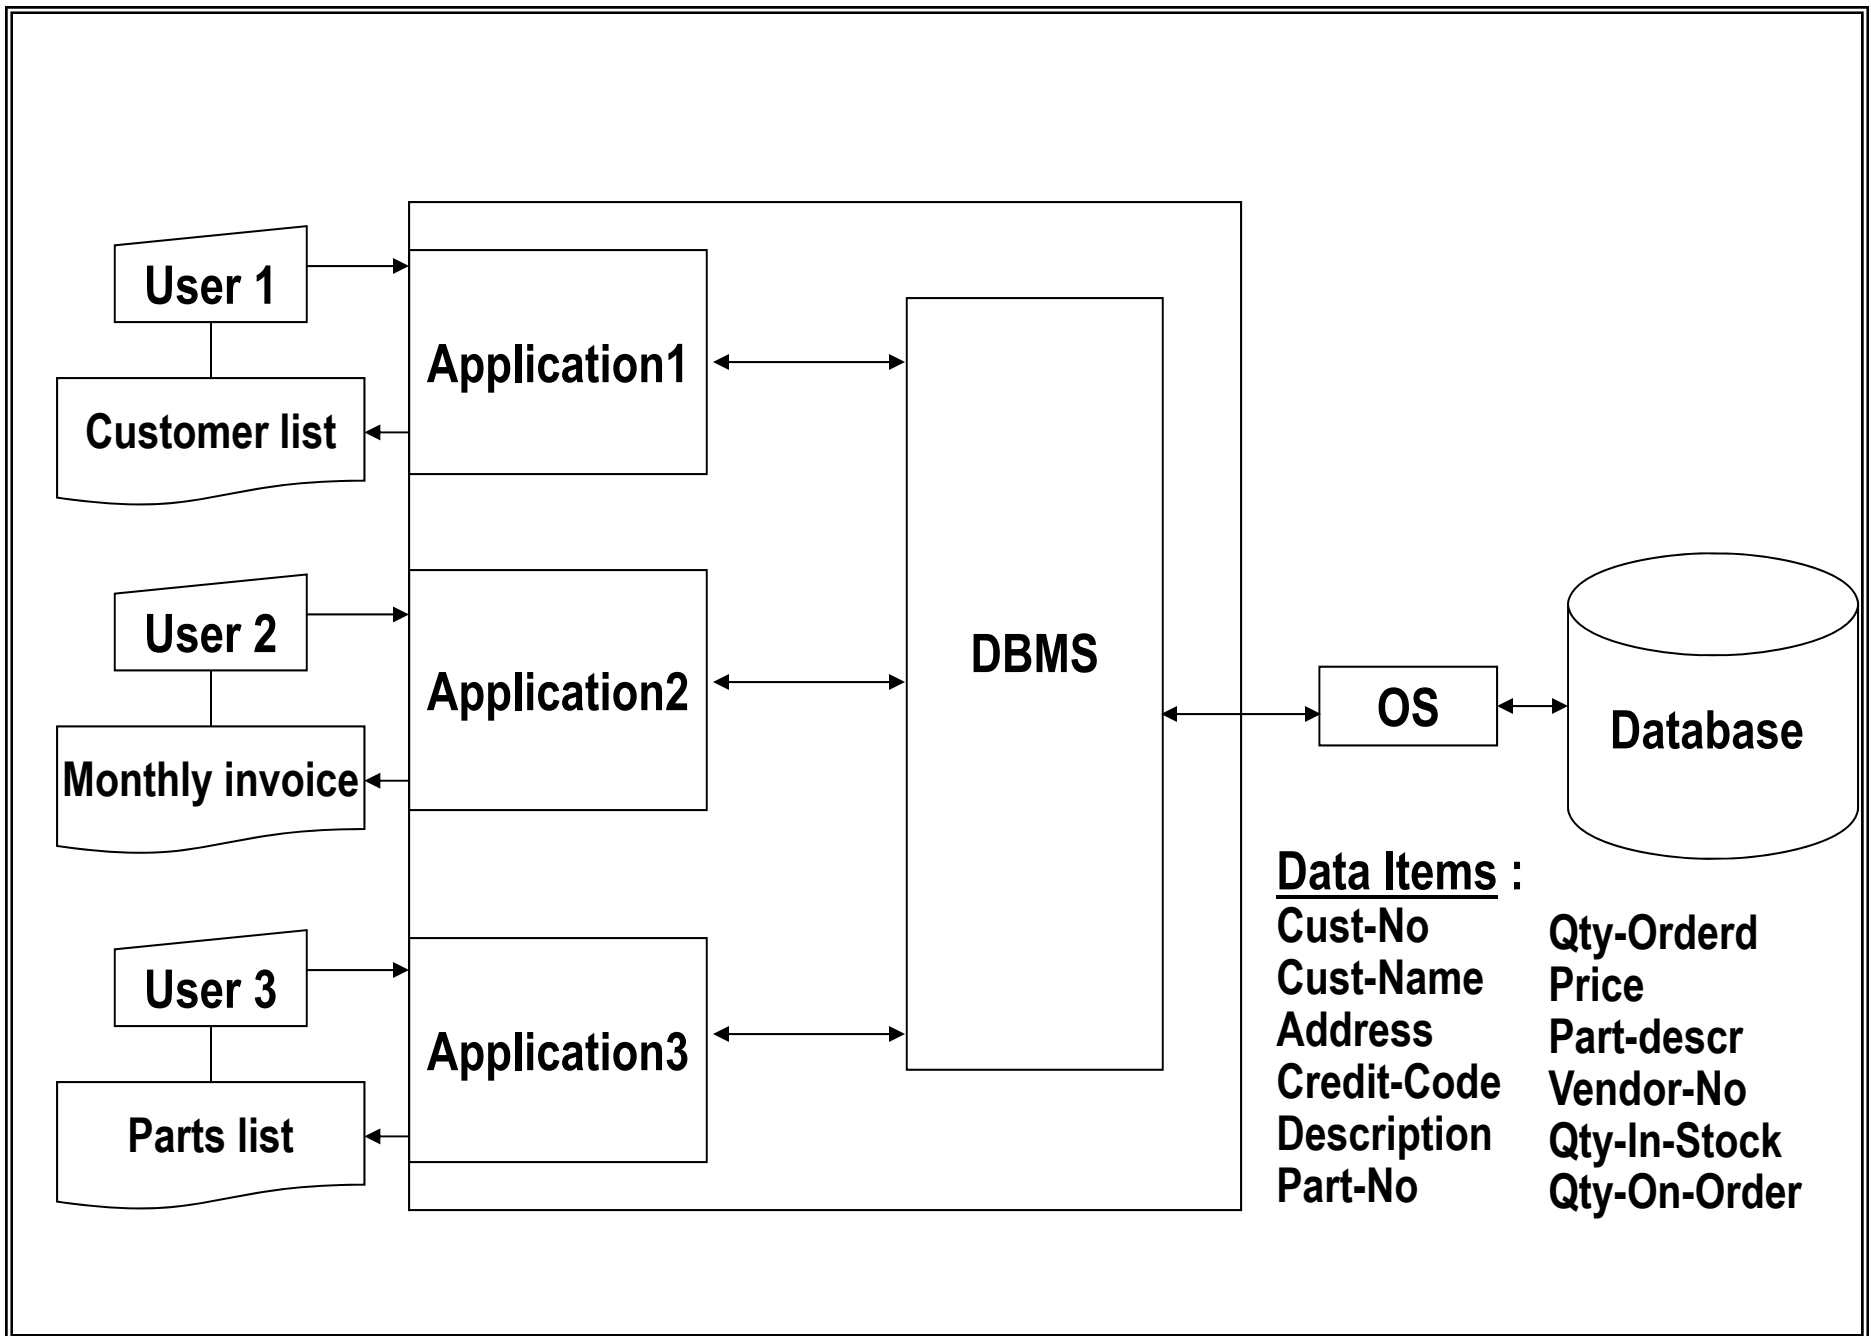
Multiple updates: Changes have to be made to all files containing the same data. *Possibility of inconsistency*

Waste of storage space:

Reliability and better local control







# Pros & Cons of DBMS

Reduce data redundancy and avoiding inconsistency

Provide Concurrent access

Offer Centralized control

- security(appropriate authorization and its control),
- integrity(constraints and their enforcements)
- reliability(backups and replication)

Data abstraction and independence

# First Step: Data Models

- ✧ Data Model: concept to describe data
- ✧ Schema: description of a collection of data using a specific data model
- ✧ Relational Model: Based on the concept of **relation** (*table with rows and columns*)

A Data Model is a collection of concepts for describing  
Entities(objects) and relationships among them  
Expressing the semantics and constraints from the real world

## Object-Based Modeling Techniques

Entity-Relationship (ER) Model

Object-Oriented (OO) Model

## Record-Based Models

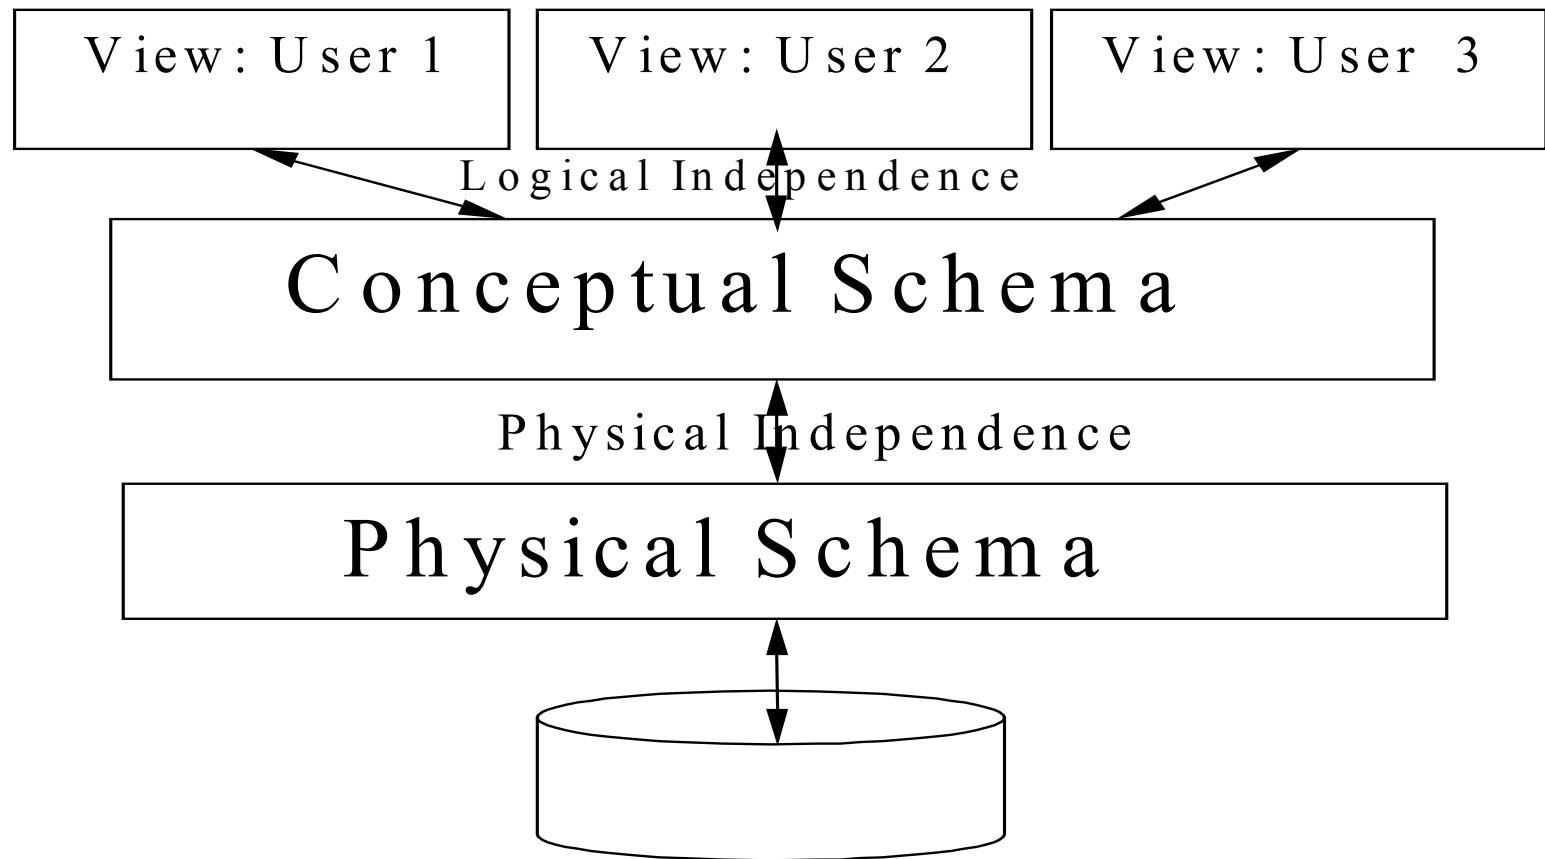
Hierarchical Model: used by earliest DBMS – IBM's IMS

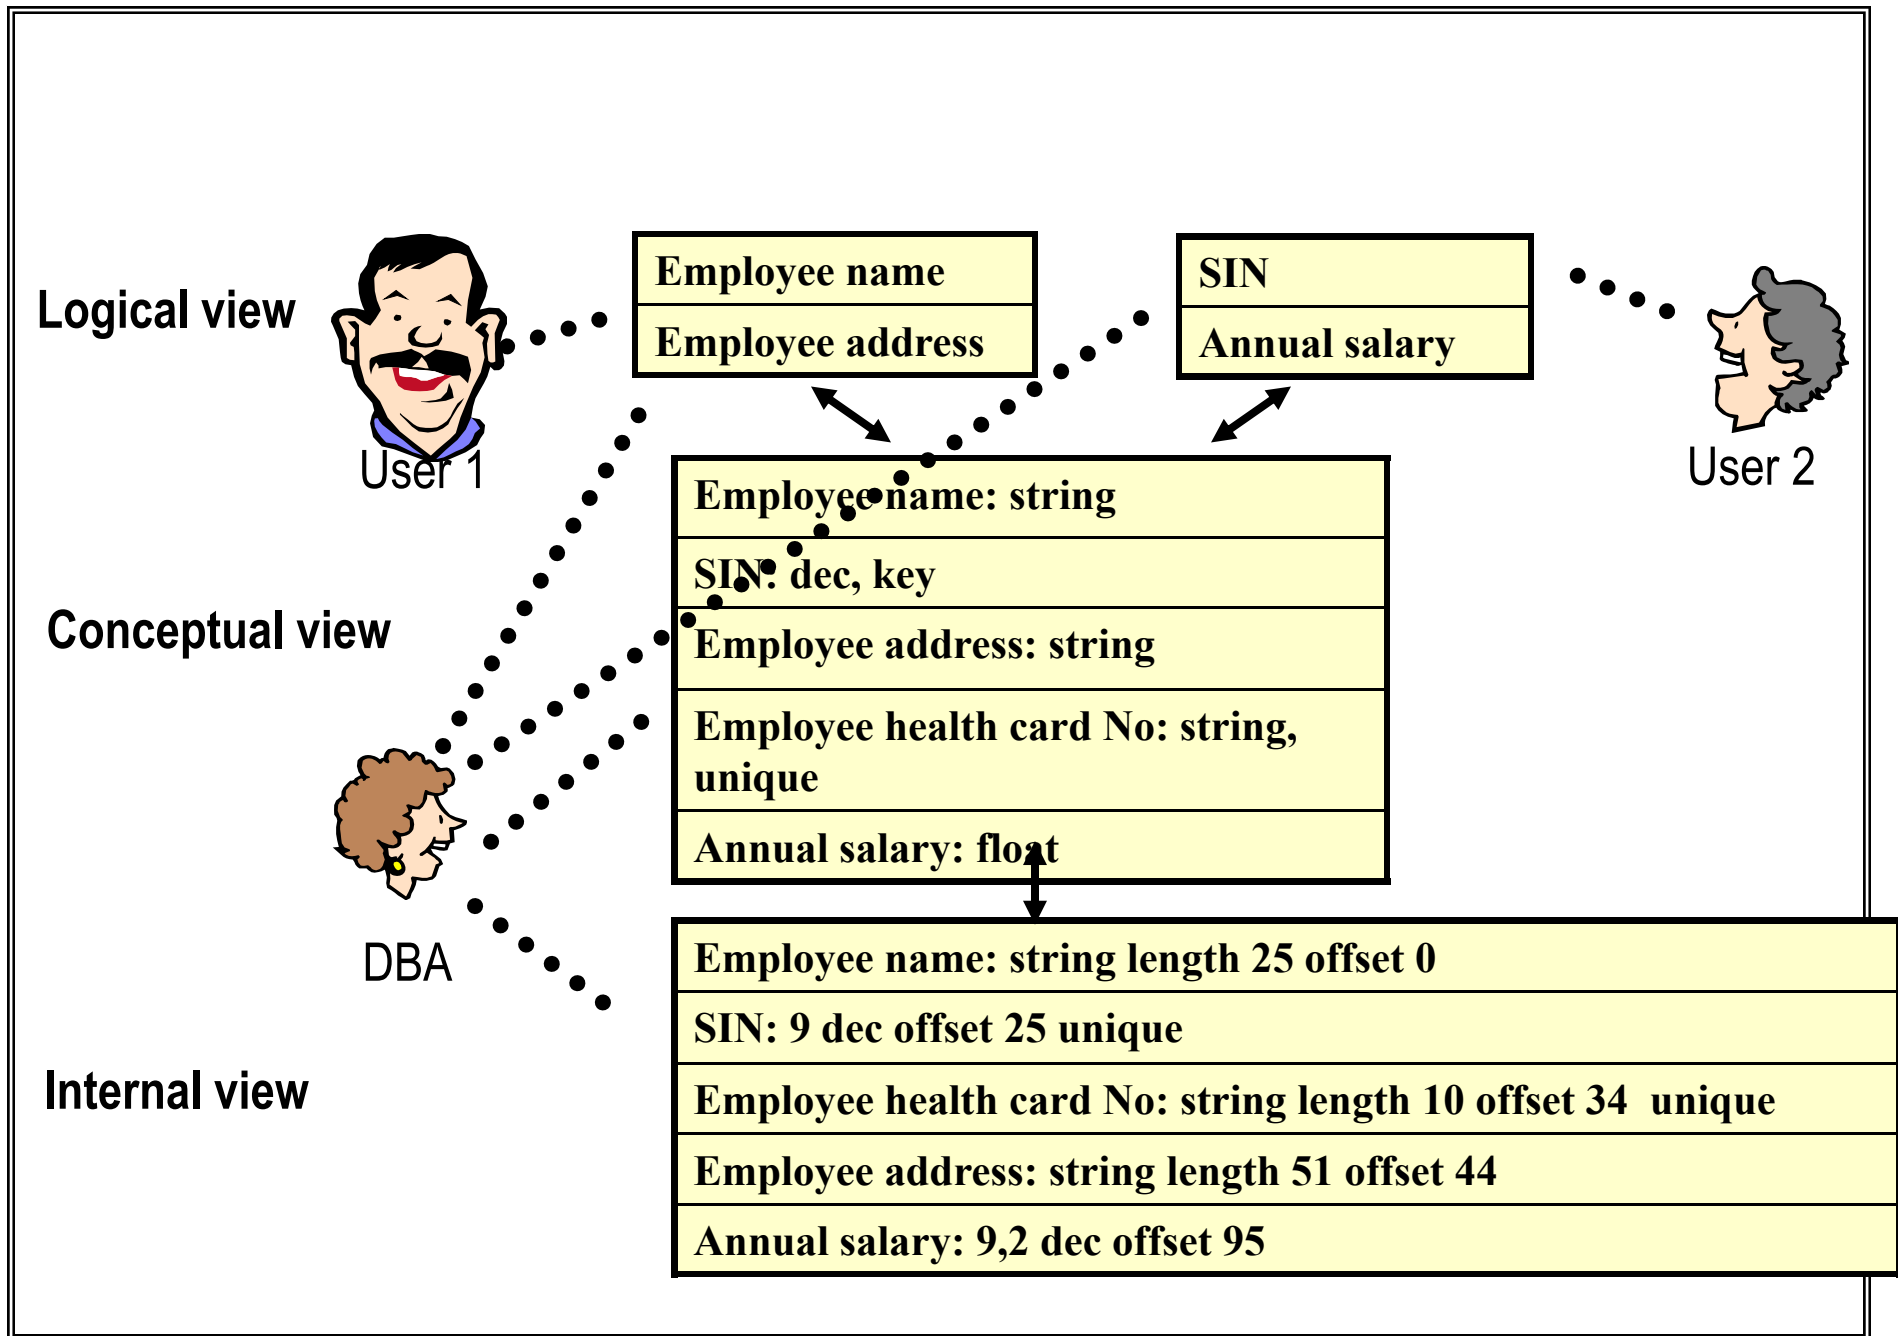
Network Model: second generation DBMS - DBTG

Relational Model: first based on theory of relations  
(RA, RC, Datalog)

Employee Name Employee Phone Number	Logical View	Employee Name Employee SIN Employee Salary
Conceptual View	Employee Name Employee Phone Number Employee SIN Employee Address Employee Annual Salary Employee YTD Salary	
Internal View	Employee Name string Employee Phone Number digits Employee SIN digits Employee Address string Employee Annual Salary moneyunits Employee YTD Salary money units	

# Three level Concepts





# Three levels & Independence

- ❖ User View: How users view data - derived from conceptual view-
- ❖ Conceptual Schema: Logical structure of the database
- ❖ Physical Schema: The actual files and indices used
- ❖ Schema defined using DDL



**Data Independence:** modify definition of schema at one level without affecting a schema definition at a higher level.

**Logical Data Independence:** modify logical schema without causing application programs to be rewritten

adding new fields to a record or changing the type of a field

**Physical Data Independence:** modify physical schema without causing logical schema or applications to be rewritten

changing file structure from sequential to direct access

# University Database

## ❖ External Schema:

Course\_Enrol(C#:char, Number:int);

## ❖ Conceptual Schema:

Student(S#, Name, Dept)

Course(C#, Cname, Credits)

Enrollment(C#, S#, grade)

## ❖ Physical Schema:

files, indexed on S#, C#, etc

A **database schema** is a description of a particular collection of data, using a given data model

Part of a schema for a university. database in relation model would contain among others, the following:

Students (sid, name, department, dob, address)

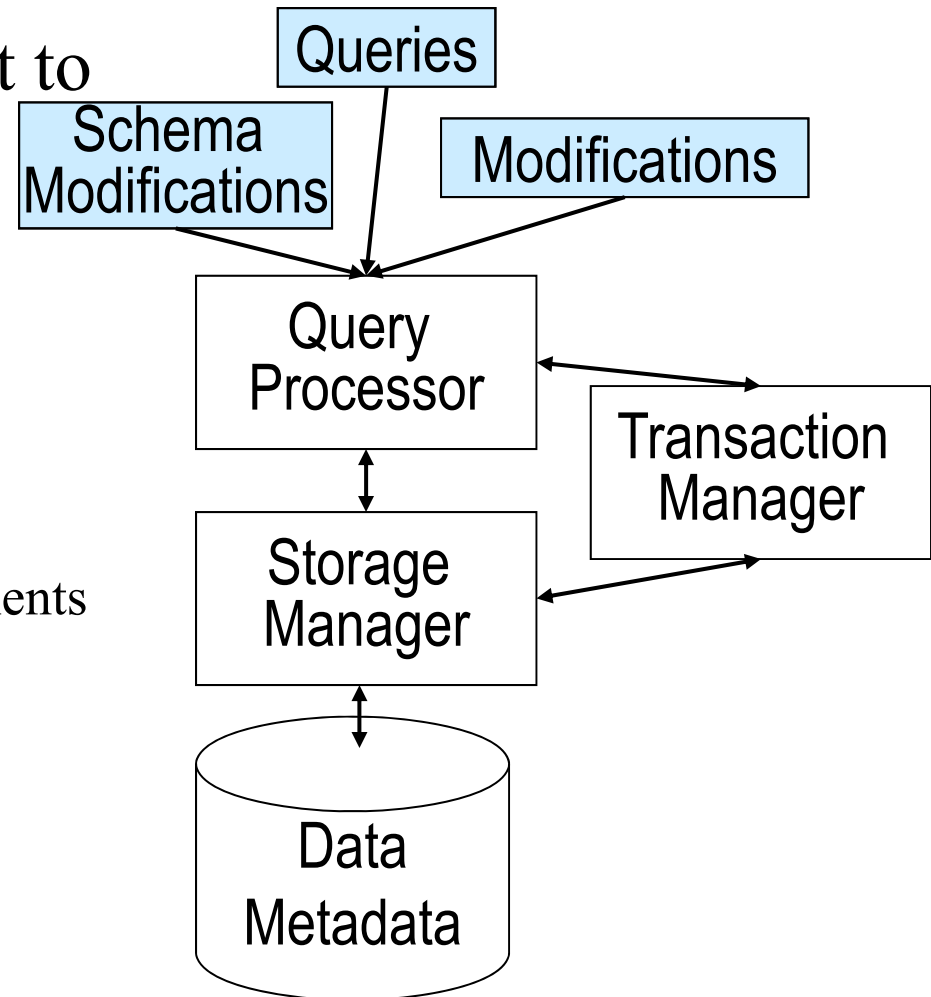
An **instance** of a database schema is the actual content of the database at a particular point in time

sid	name	department	dob	address
1112223	John Smith	CS	12-01-82	22 Pine, #1203
2223334	Ali Brown	EE	31-08-73	2000 St. Marc
3334445	Youwong Li	CS	23-11-79	1150 Guy

# The Architecture of a DBMS

❖ There are 3 types of input to DBMS:

- ◆ Access via queries
- ◆ Updates to data
- ◆ Updates to model
  - ✧ Initial database creation,
  - ✧ addition to schema components
  - ✧ schema modifications

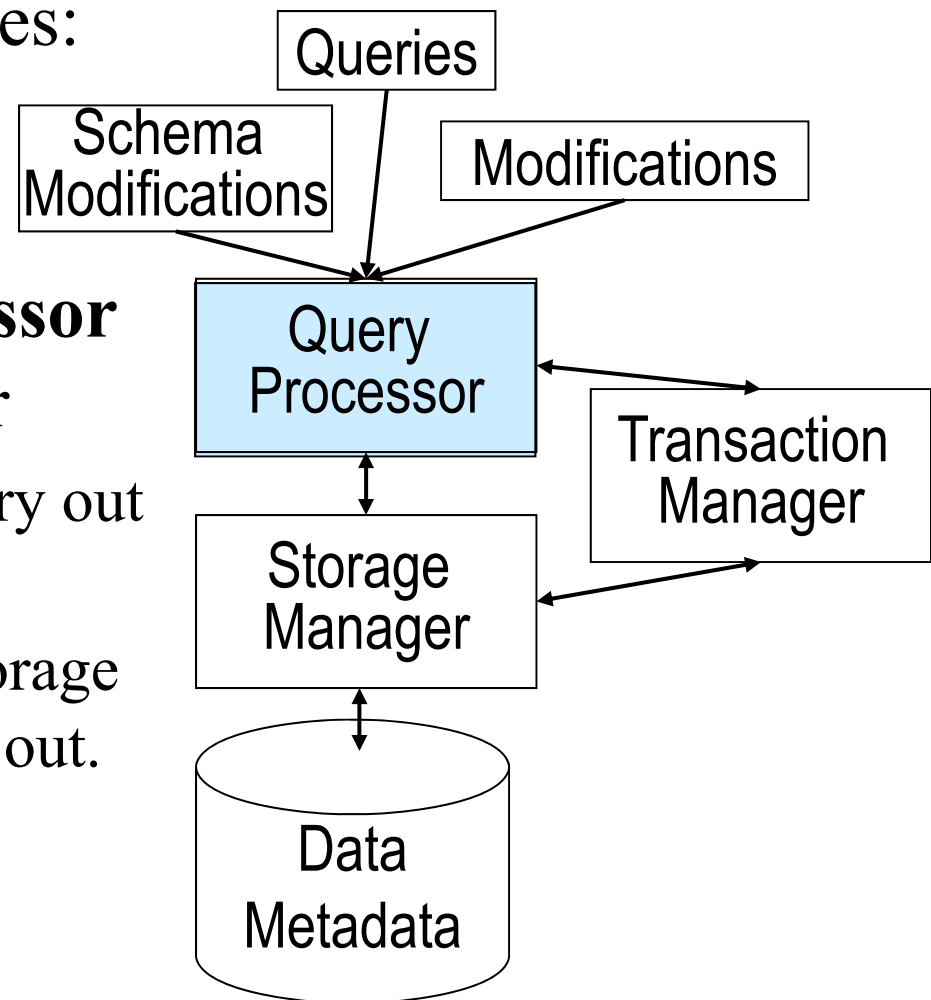


❖ The **query processor** handles:

- ◆ Queries
- ◆ Updates

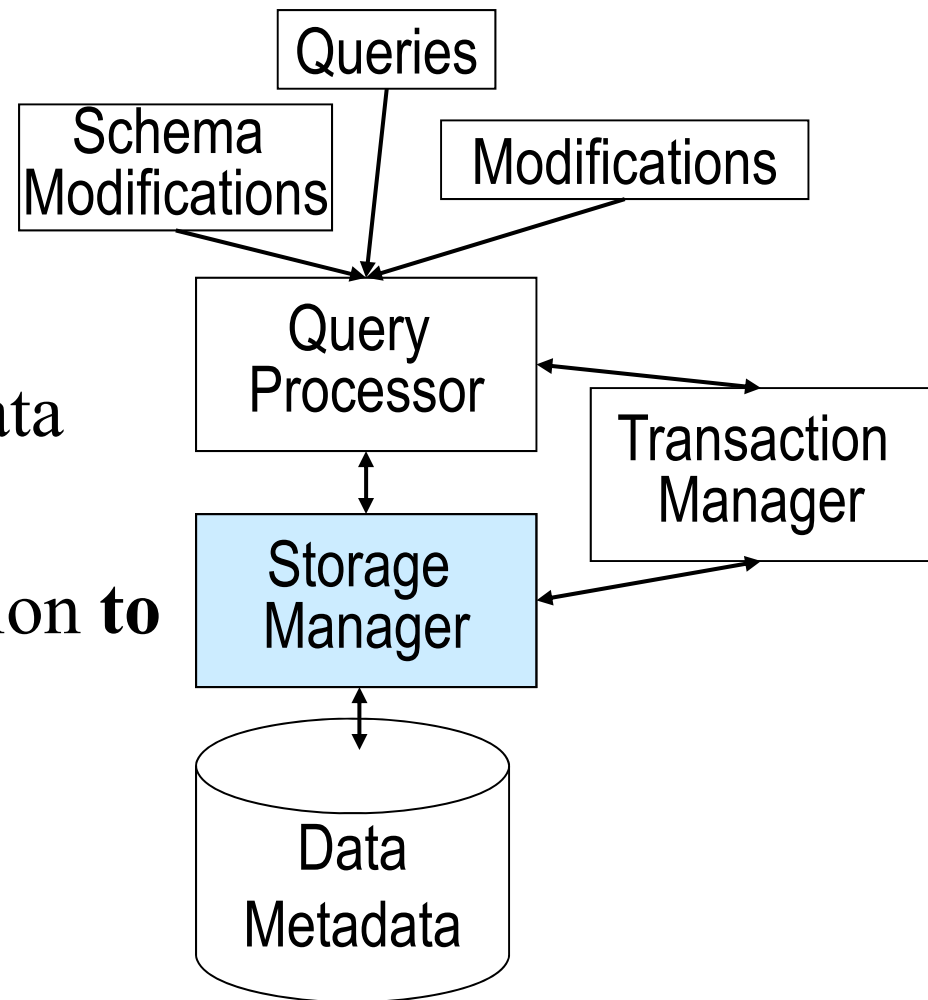
❖ The job of the **query processor** which includes an optimizer

- ◆ To find the “best” way to carry out a requested operation
- ◆ To issue commands to the storage manager that will carry them out.



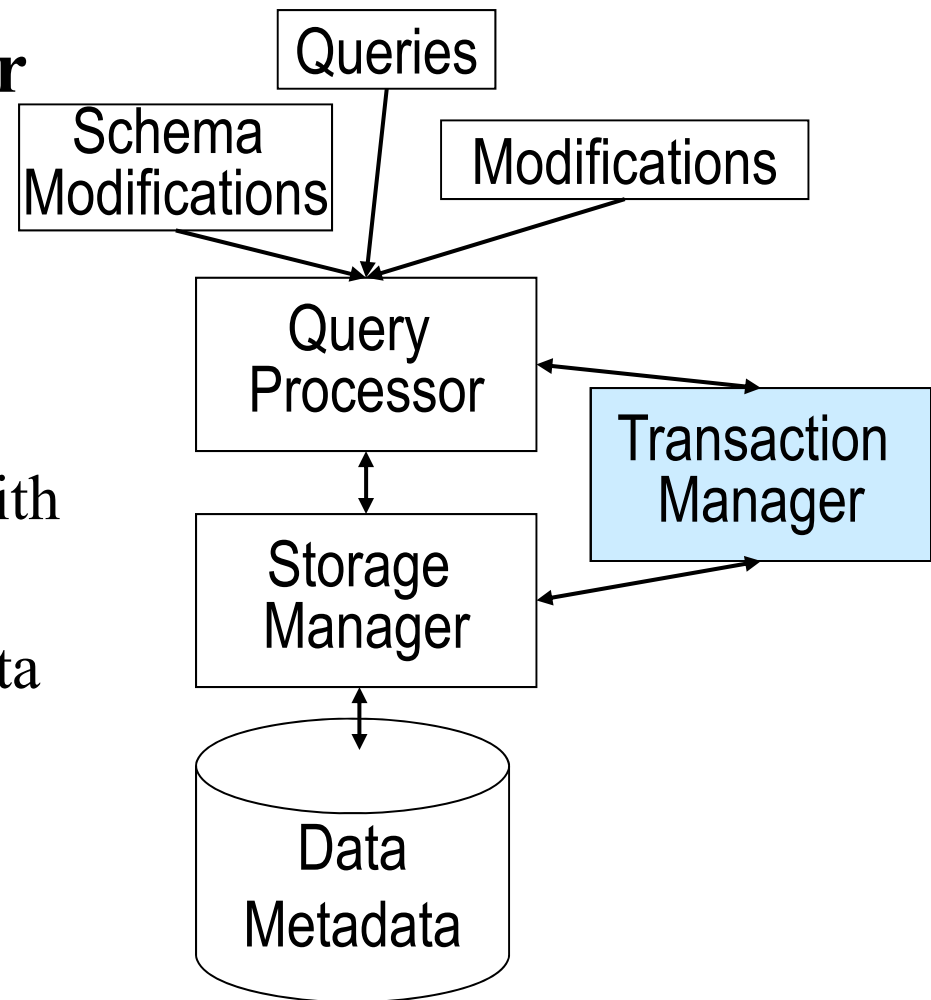
❖ The job of the **storage manager** is

- ◆ To obtain requested information **from** the data storage
- ◆ To modify the information **to** the data storage when requested.



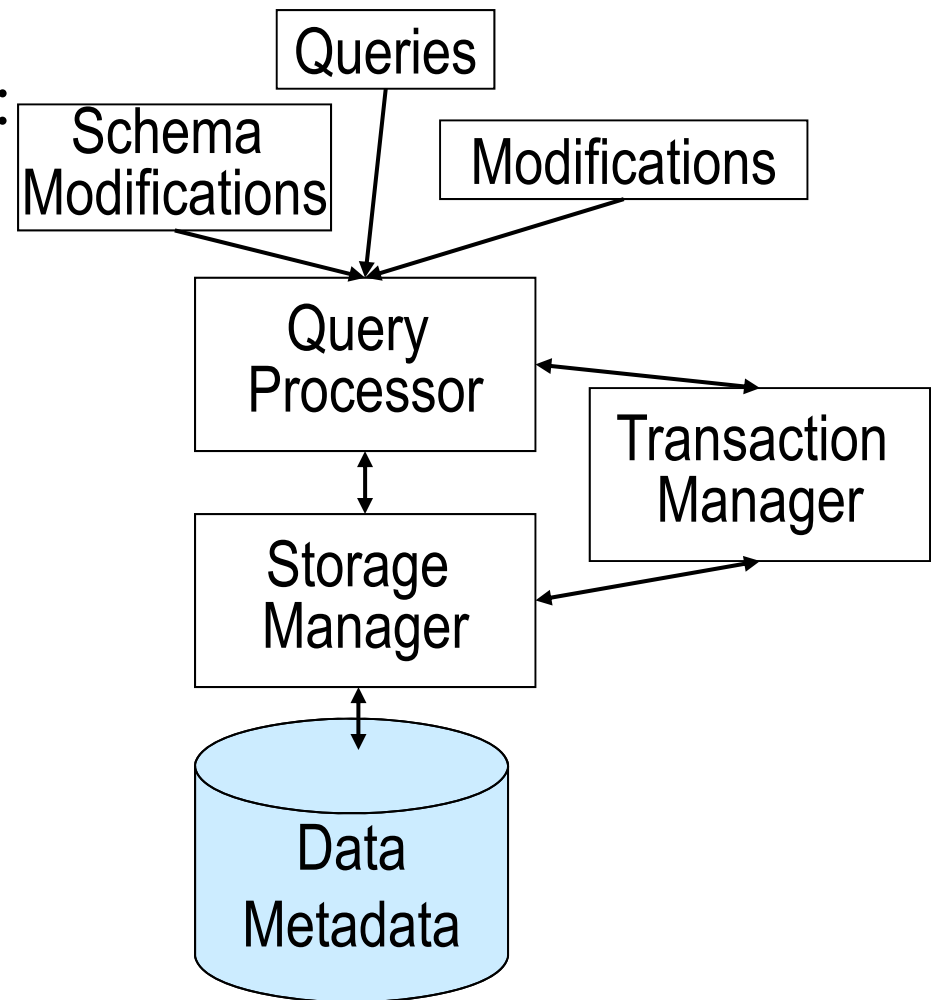
❖ **The transaction manager**  
is responsible for the  
**enforcing ACIDity**

- ◆ several concurrent transactions (one or more queries) do not interfere with each other
- ◆ the system will not lose data even if there are failures (*done through Recovery subsystem*)



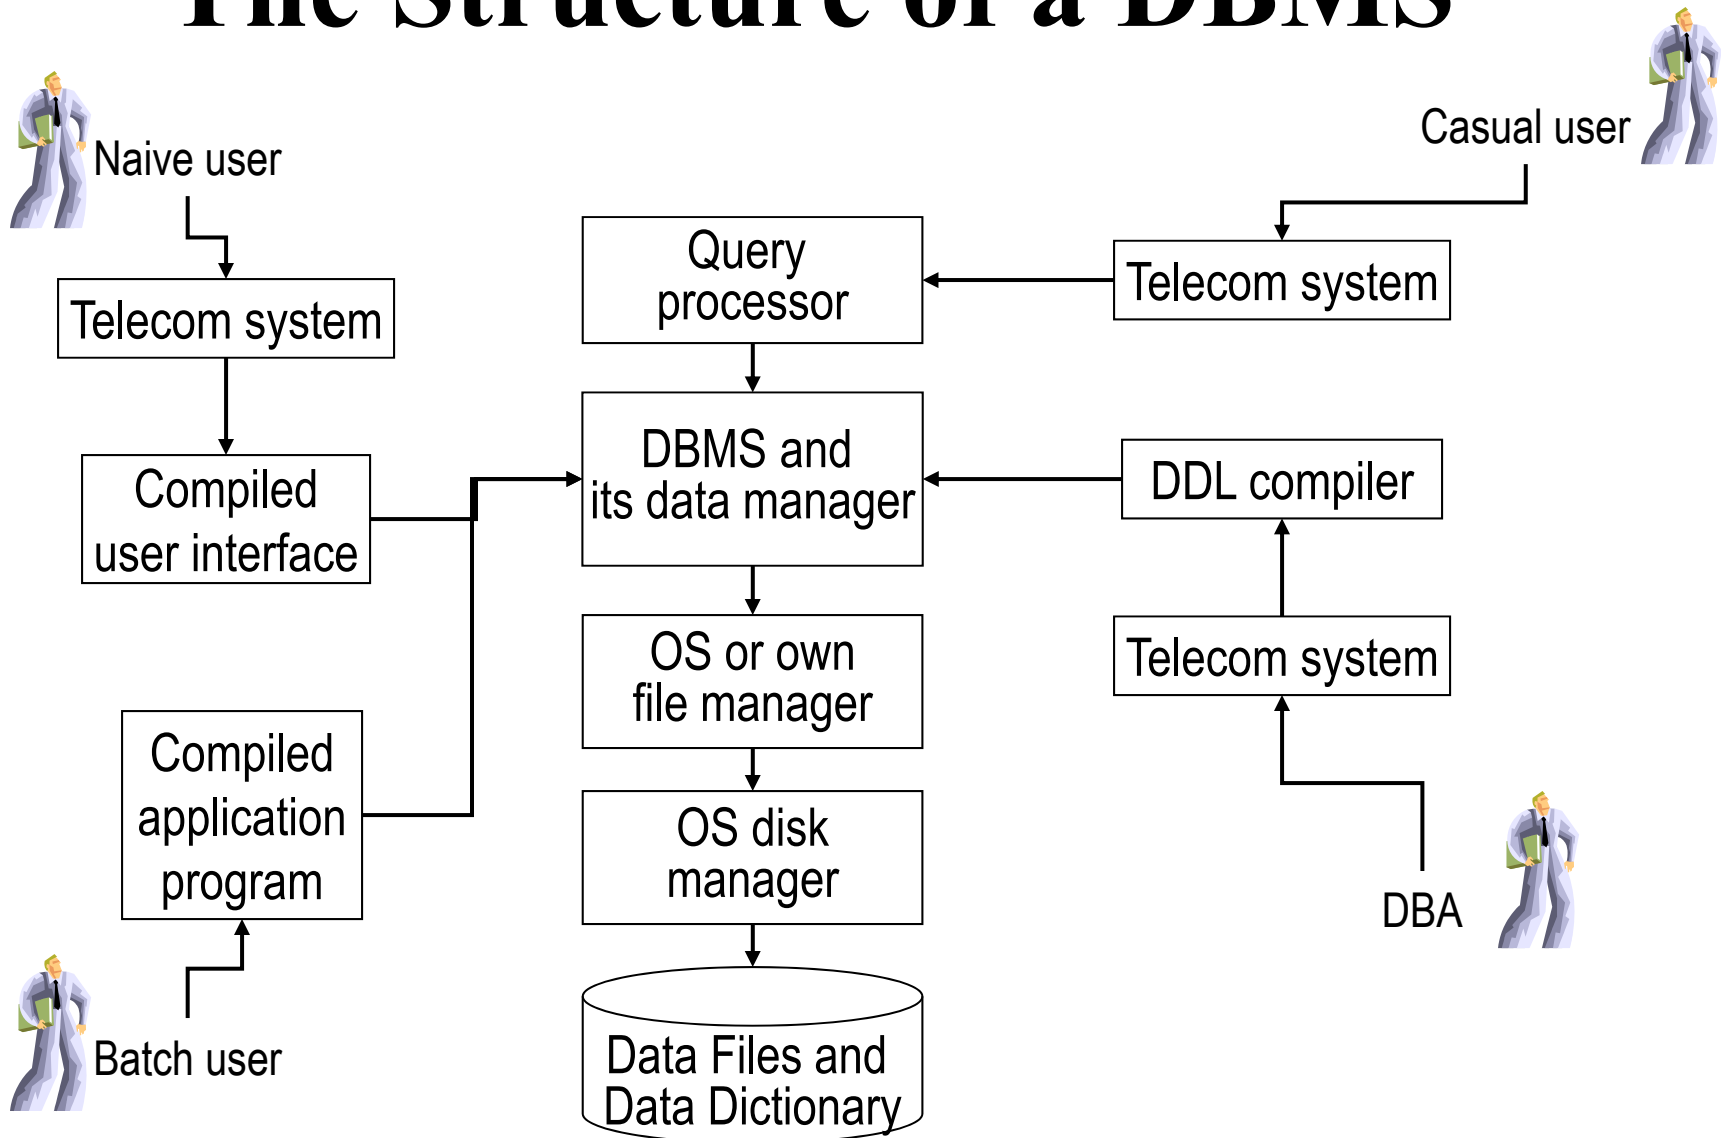
❖ Database contents include:

- ◆ Metadata for the DBMS and one or more databases
- ◆ Data belonging to one or more databases
- ◆ Access aids such as indices and statistics

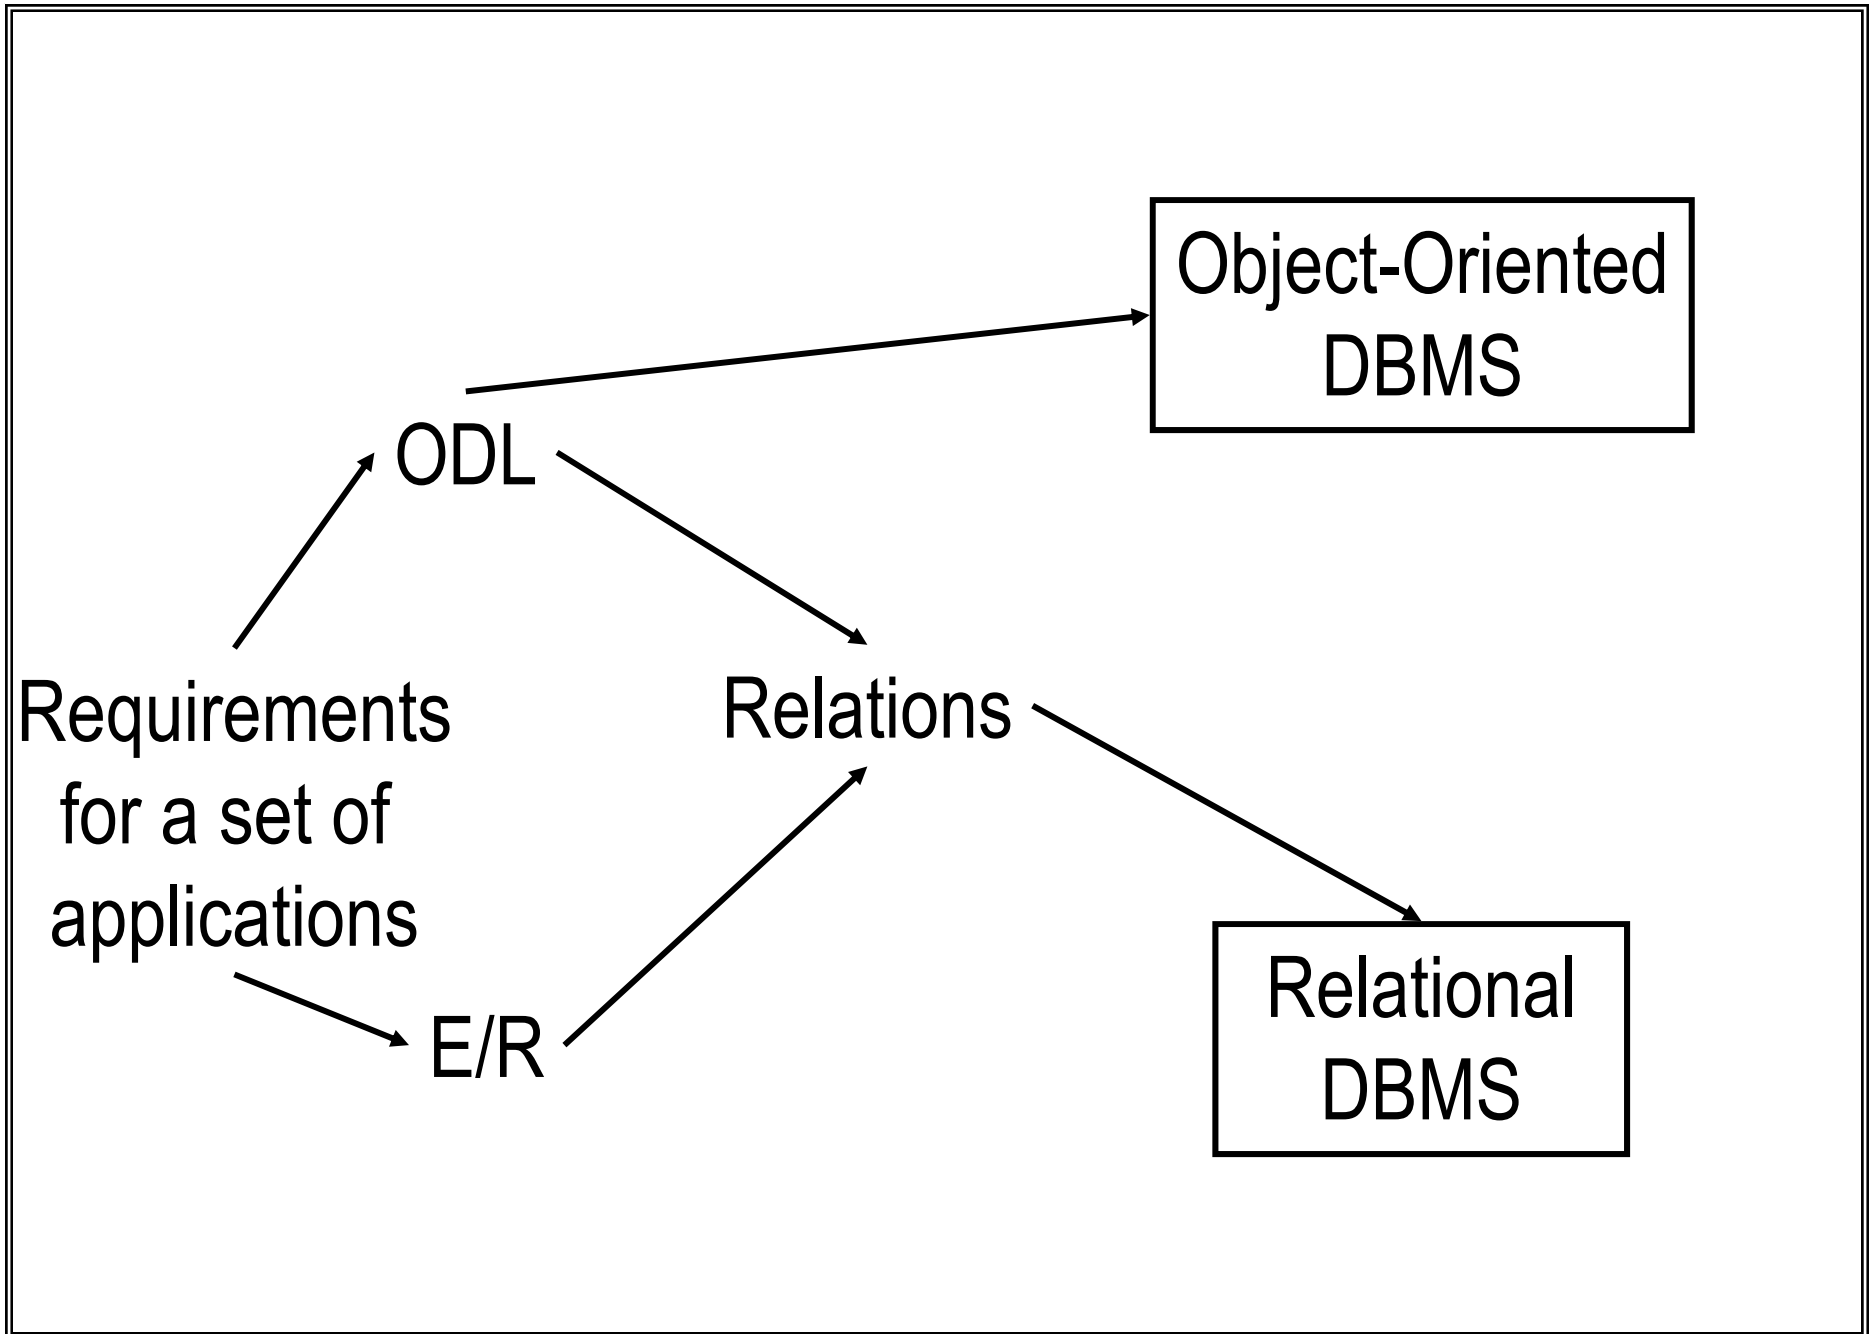


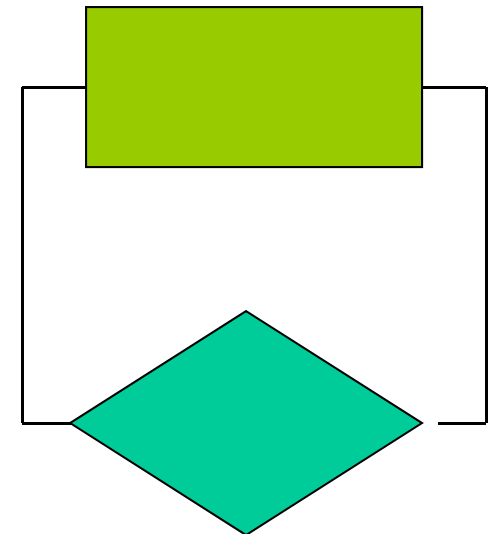
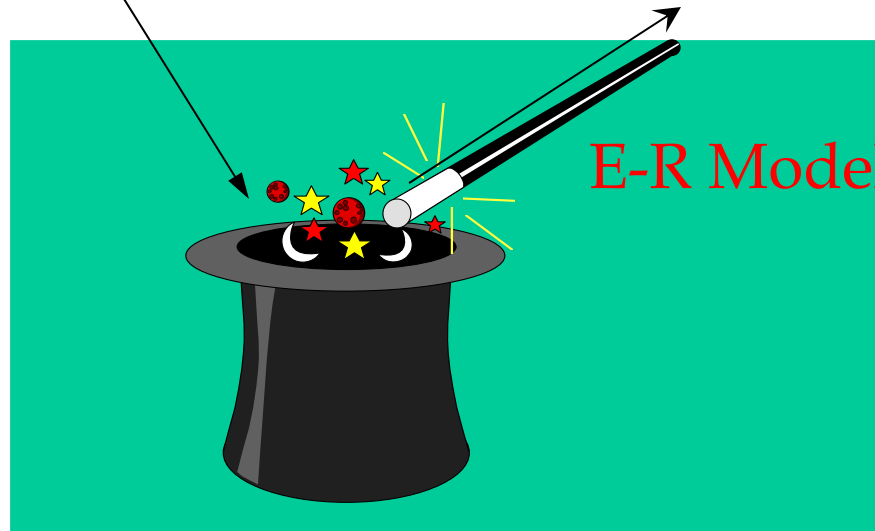
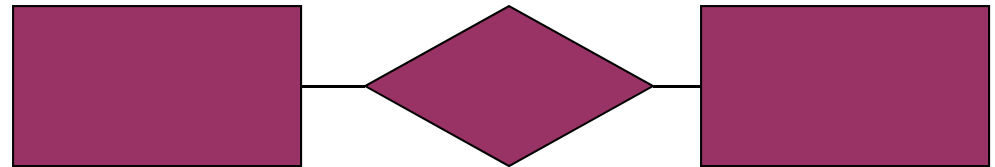
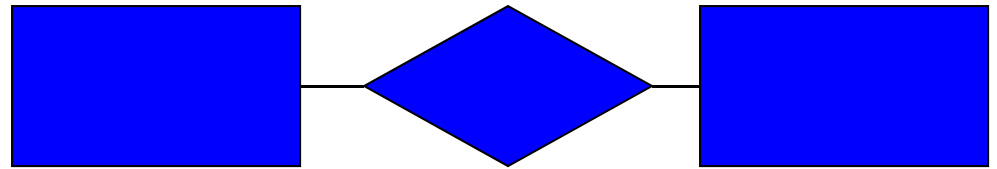
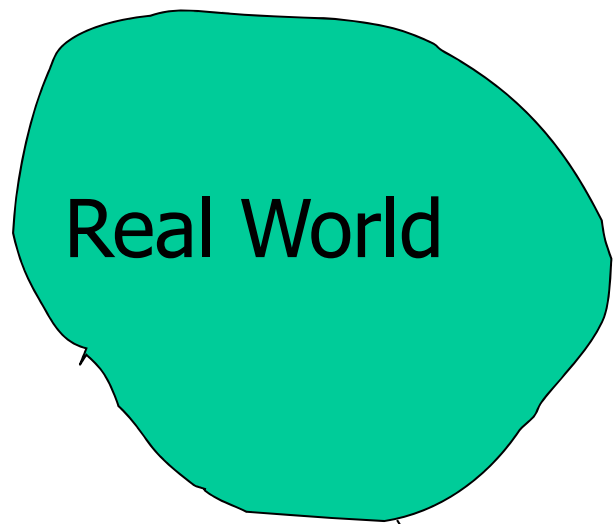


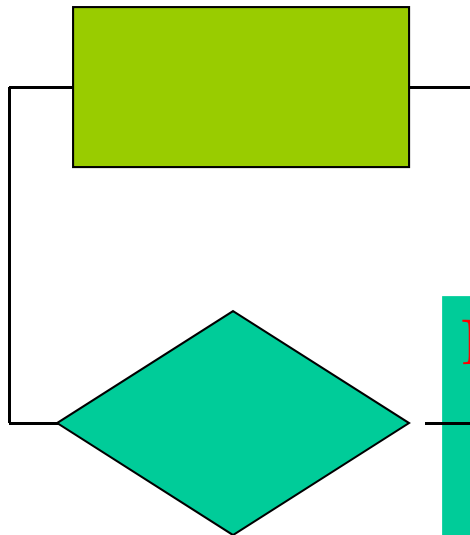
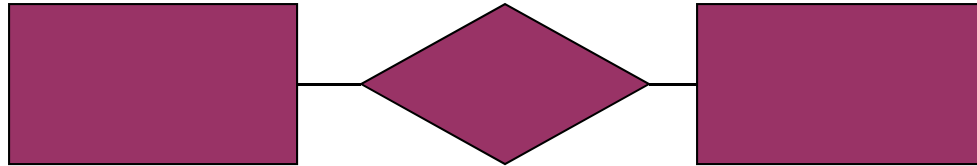
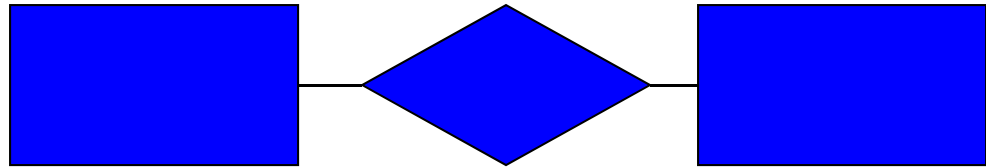
# The Structure of a DBMS



# **Database Design Process and Conceptual Design**







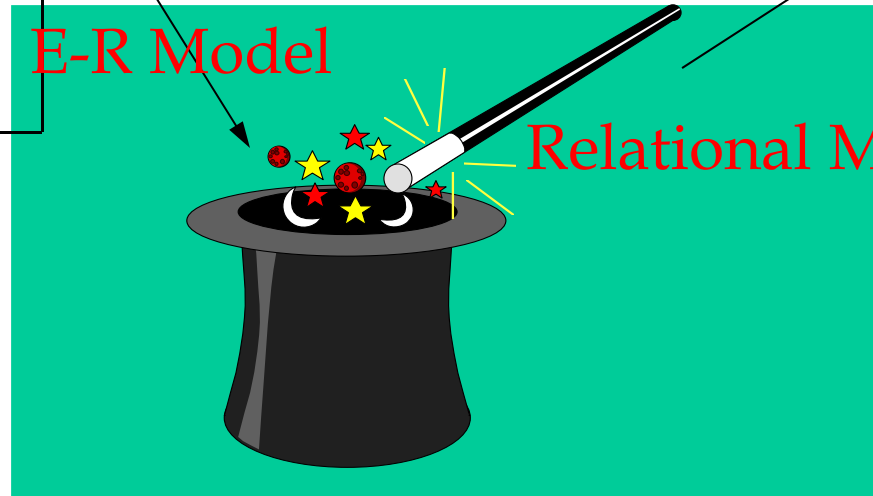
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc

•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc

•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc
•abc	Abc 123q 456	•abc

E-R Model

Relational Model



# Relational Model

In this model, the data is organized in relations (tables)

Relational database schema: **DDL component of SQL**

- set of table names

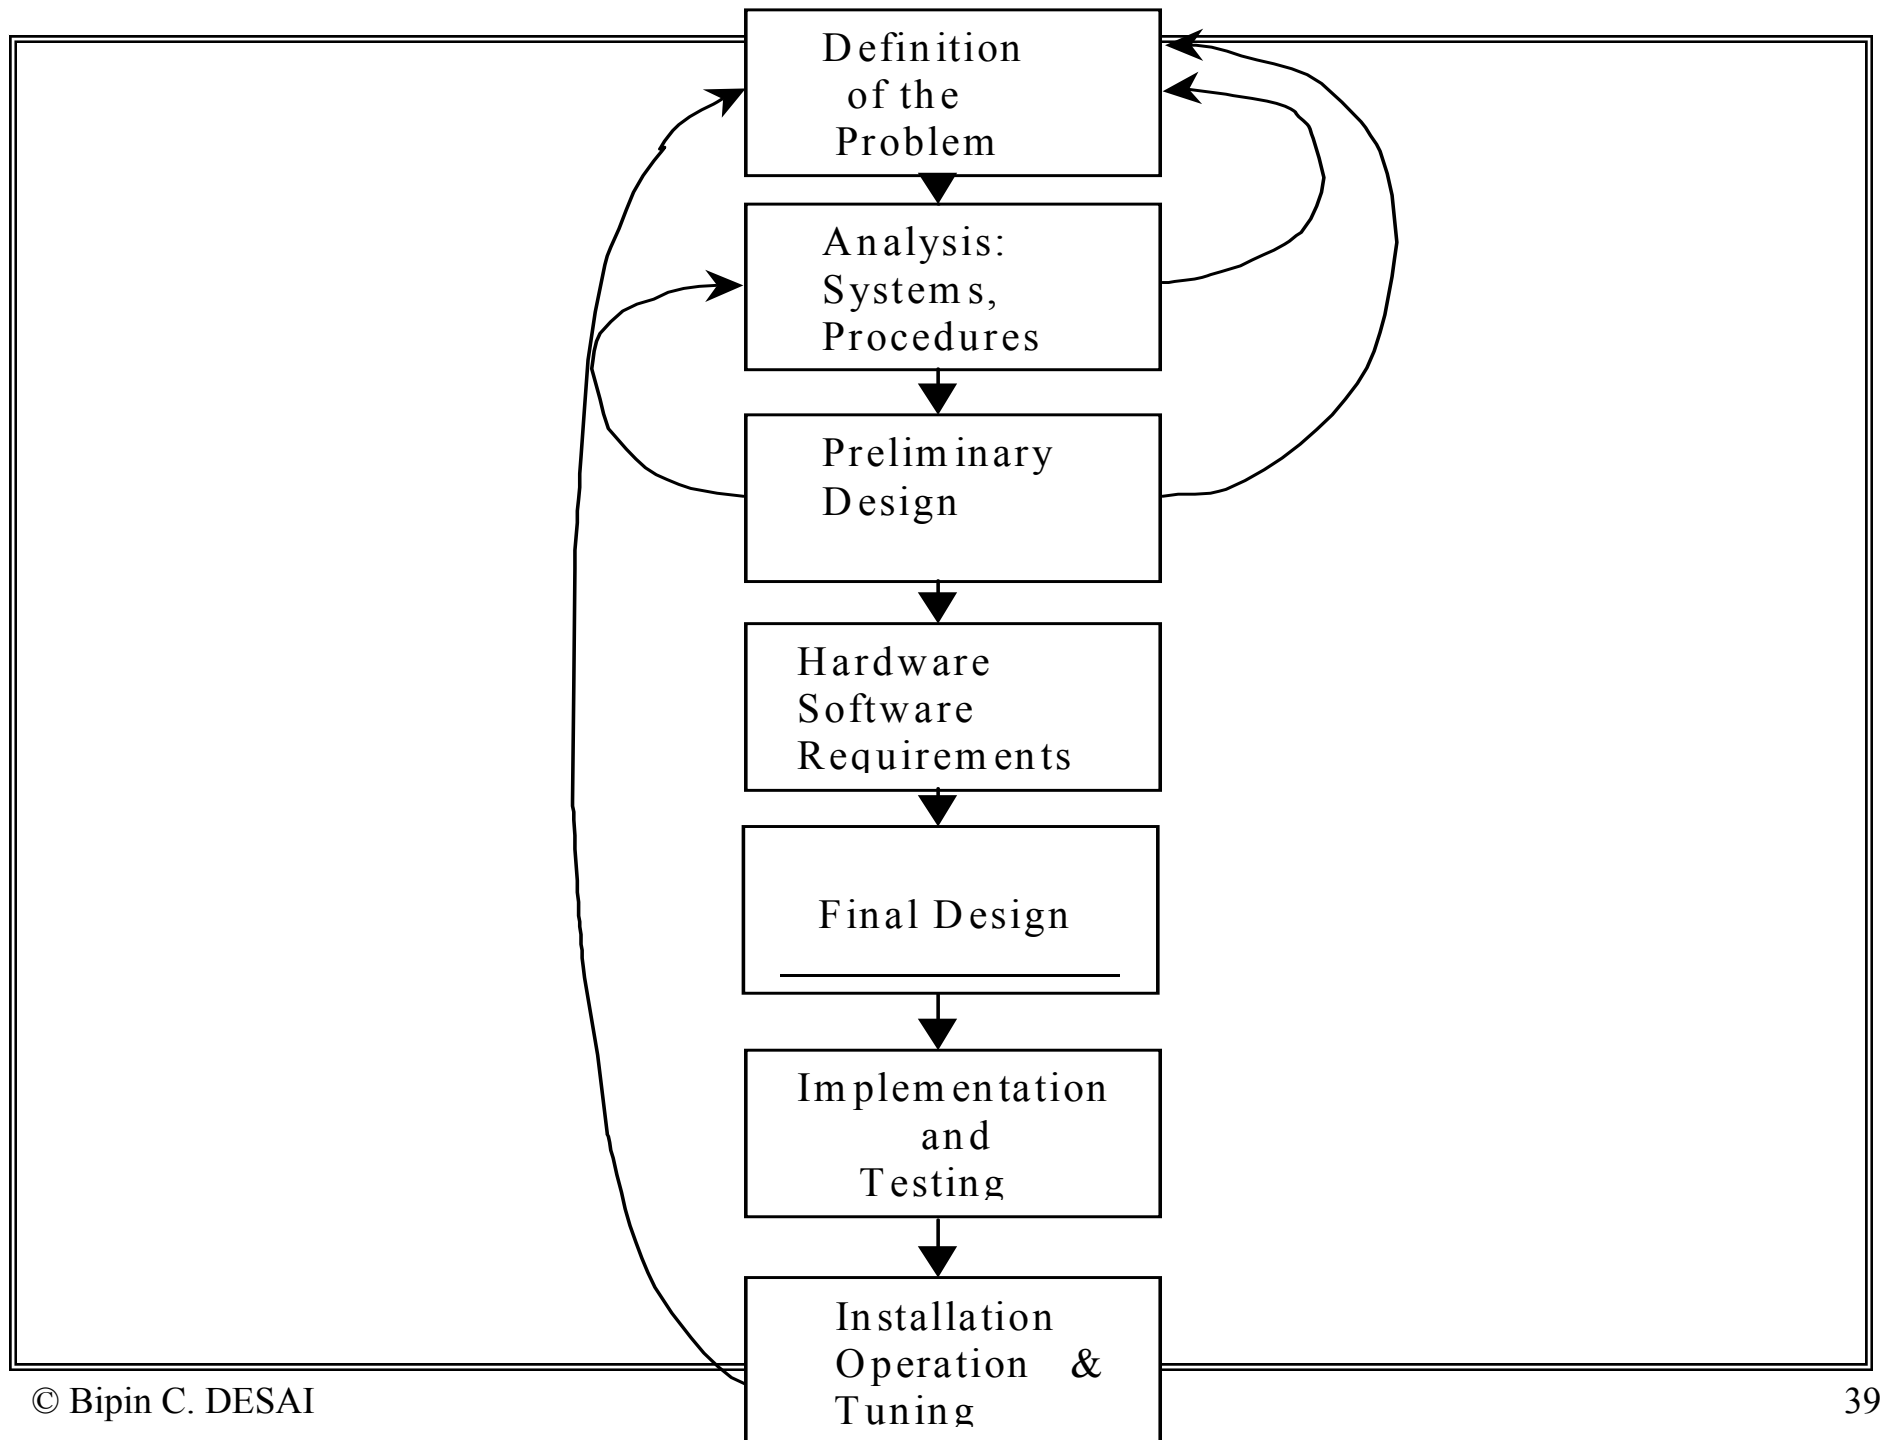
- list of attributes for each table and their properties

Examples of tables from a university database:

- Student** : stud\_number, name, address, program

- Department**: name, budget\_code, room, phone

- Course**: name, number, credits



# Database Design Process

- Definition of the problem
- Study underlying applications (*Procedure Manuals, Interviews etc.*)
  - ✦ What are the entities and relationships involved?
  - ✦ What details about them should be in the database?
  - ✦ What are the *procedures, business rules, constraints*?
  - ✦ Who are the users? What do they need?
- Preliminary Conceptual design:
  - ✦ ER Model



# Database Design Process

- Software/Hardware Requirements

- ✦ UML for software design

- Final Design: Schema Refinement: (Normalization)

- ✦ Check relational schema for redundancies and related anomalies.

- ✦ External Schemas, indices, views, access methods

- Application programs, forms, reports, user interfaces

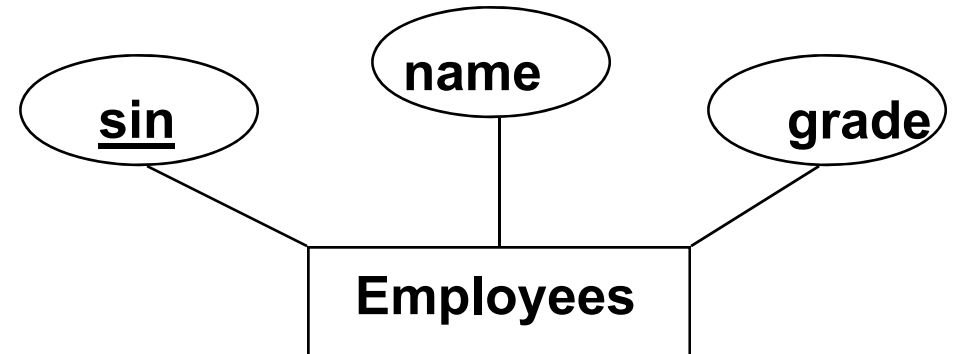
- Implementation and testing

- Installation and Tuning:

- ✦ Data Distribution, Physical re-design

- ✦ Performance, Security, Backup & Recovery.

# ER Model



❖ Entity: Real-world object distinguishable from other objects.

◆ An entity is described using a set of attributes.

❖ Entity Set: A collection of similar entities.

◆ All entities in an entity set have the same set of attributes.

◆ Each entity set has a key.

◆ Each attribute has a domain.

◆ Can map entity set to a relation easily.

```
CREATE TABLE Employees
(sin CHAR(9),
name CHAR(25),
grade INTEGER,
PRIMARY KEY (sin))
```

```
mysql> CREATE TABLE Employees
        (sin CHAR(9),
         name CHAR(25),
         grade INTEGER,
         PRIMARY KEY (sin));
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> show tables;
+-----+
| Tables_in_db11s |
+-----+
| Employees       |
+-----+
```

```
mysql> desc Employees;
```

Field	Type	Null	Key	Default	Extra
sin	char(9)	NO	PRI		
name	char(25)	YES		NULL	
grade	int(11)	YES		NULL	

3 rows in set (0.00 sec)

The Extra field contains any additional information that is available about a given column.

The value is `auto_increment` for columns that have the `AUTO_INCREMENT` attribute and empty otherwise.

```
CREATE TABLE Department
    (did mediumint not null auto_increment,
     dname CHAR(16),
     bcode char(12),
     PRIMARY KEY (did));
```

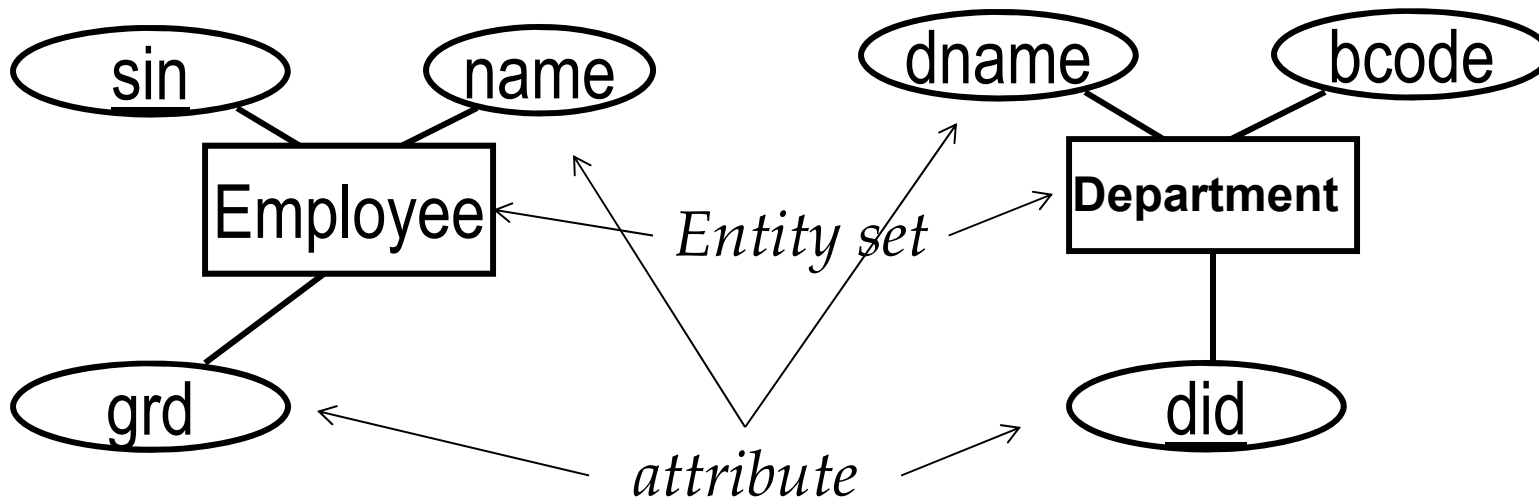
Query OK, 0 rows affected (0.04 sec)

```
mysql> desc Department;
```

Field	Type	Null	Key	Default	Extra
did	mediumint(9)	NO	PRI	NULL	auto_increment
dname	char(16)	YES		NULL	
bcode	char(12)	YES		NULL	

3 rows in set (0.00 sec)

## Entities and entity sets



All employees, and departments have the same set of properties(attributes)

To distinguish one instance of an entity in an entity set from others, we introduce an identifying attribute

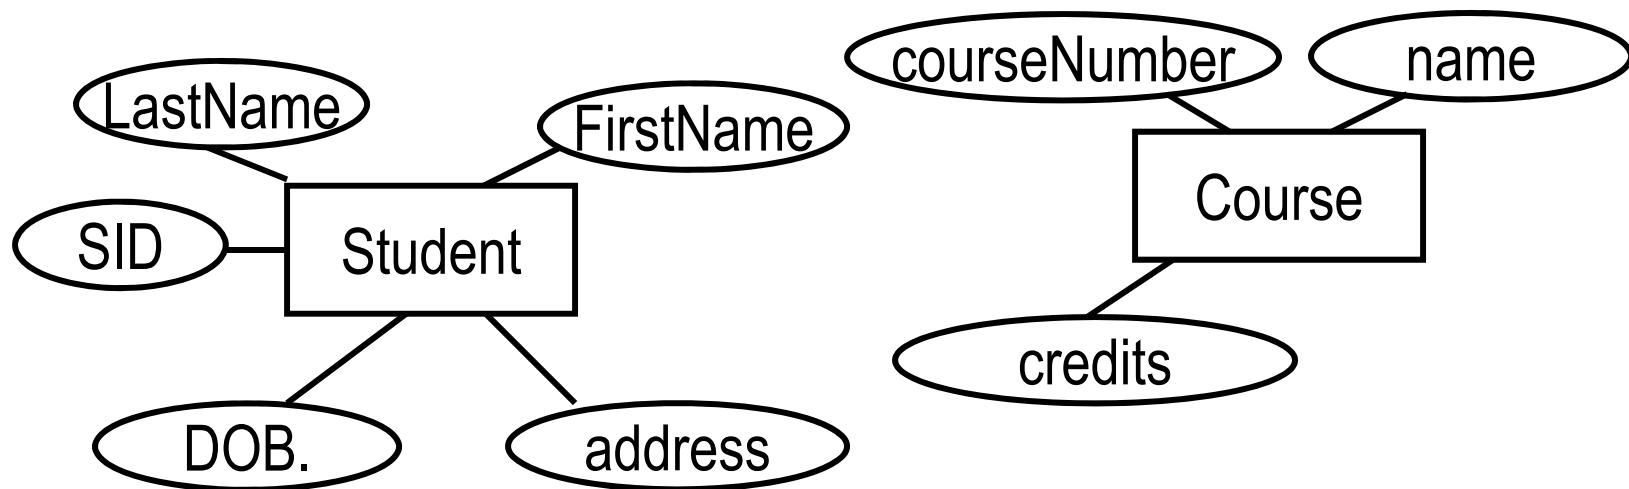
This is the primary key and it is underlined

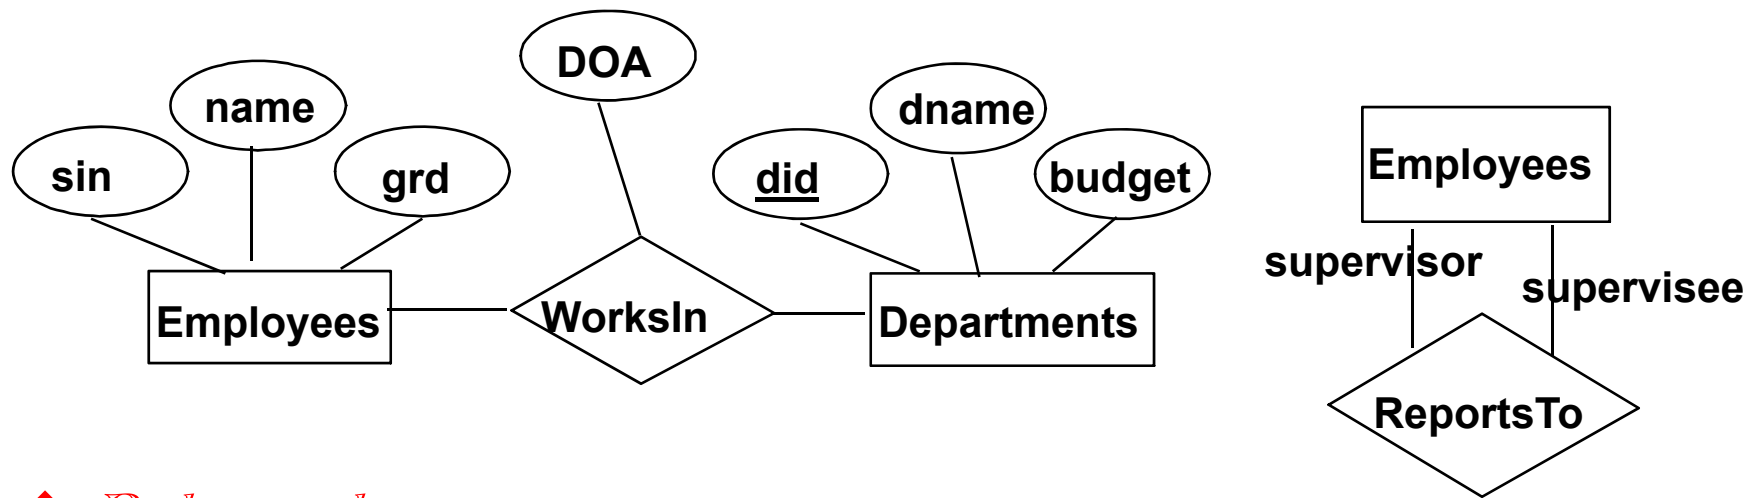
**Entity** – Real world object distinguishable from other objects of the same type

**Entity Set** -- A collection of similar entities: all have same set of properties

ODL:

**Object** corresponds to entity   **Class** corresponds to entity set





❖ Relationship:

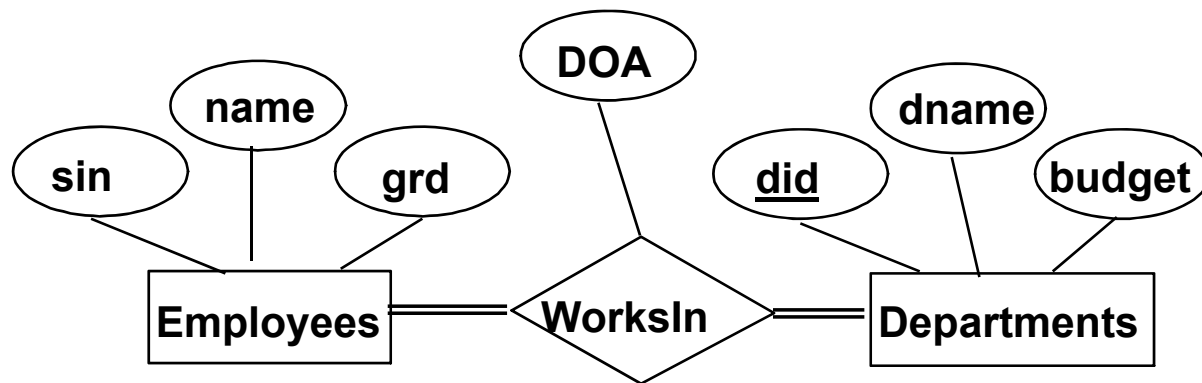
- ◆ Association among 2 or more entities.

❖ Relationship Set: Collection of similar relationships.

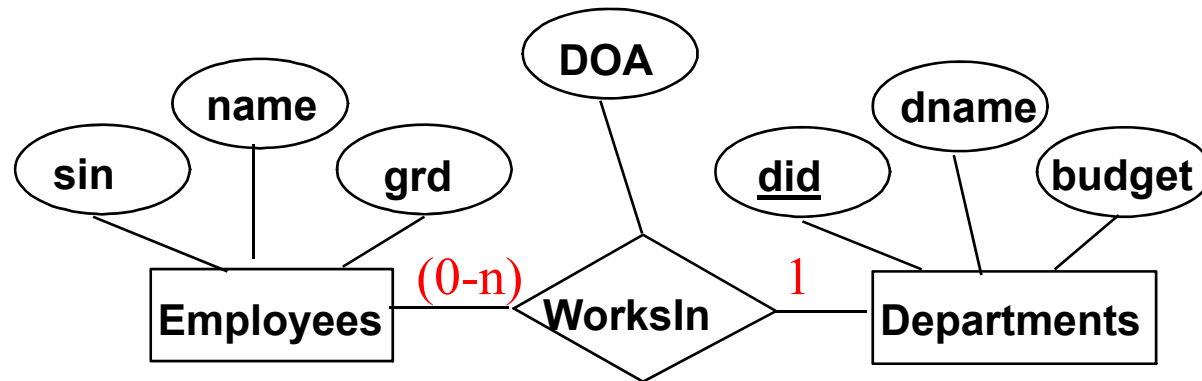
- ◆ An n-ary relationship set  $R$  expresses an association among  $n$  entity sets  $E_1 \dots E_n$ ; each relationship in  $R$  involves entities  $e_1 \in E_1, \dots, e_n \in E_n$

*Same entity set could participate in different relationship sets, or in different “roles” in same set.*

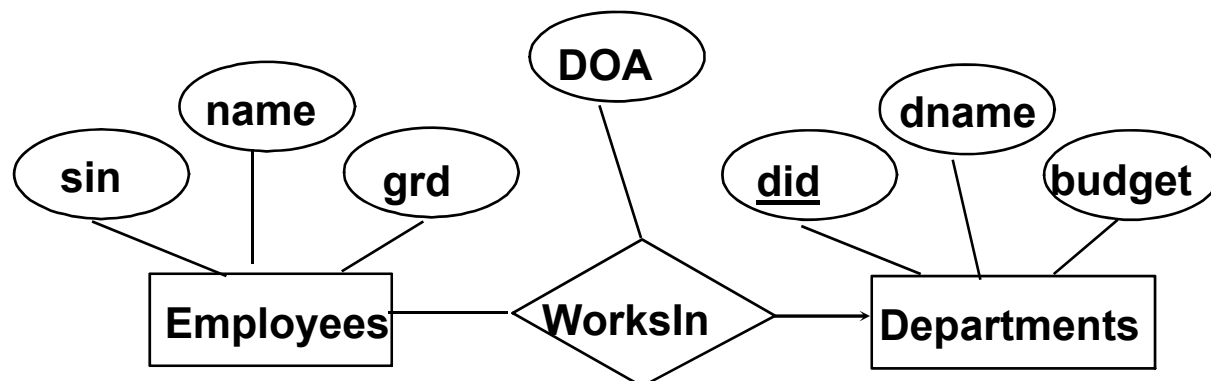




Total participation of all employees & departments  
In the WorksIn relationship

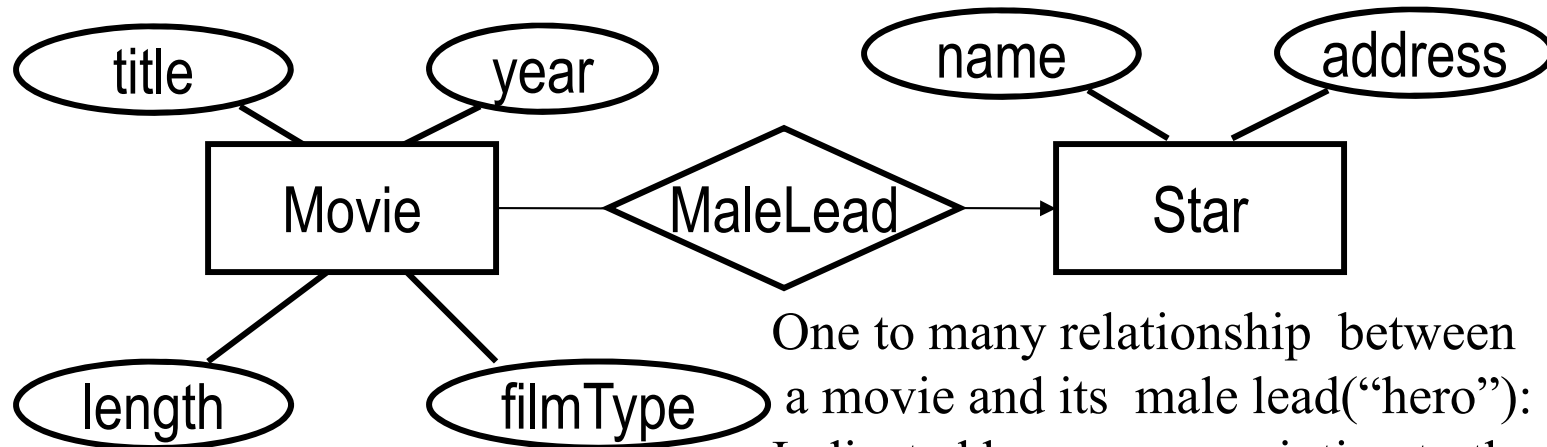
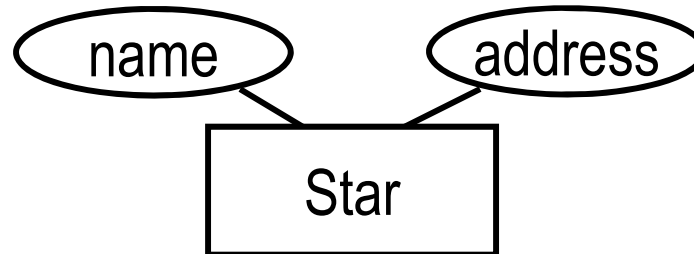
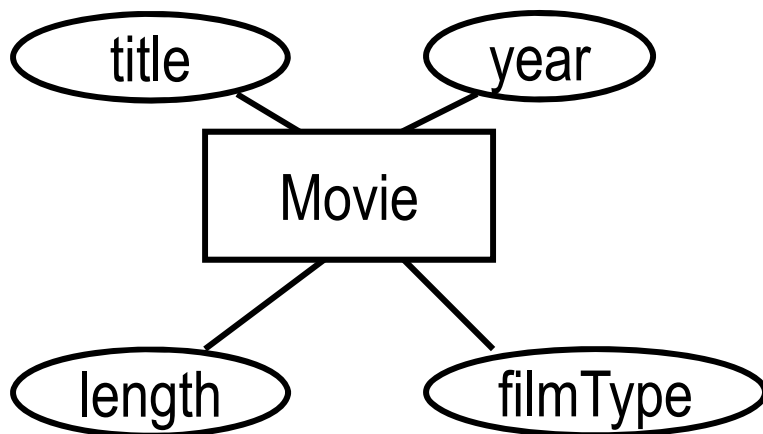


Many to one relationship:  
many employees in a  
department;  
but an employee is  
assigned to only one  
department



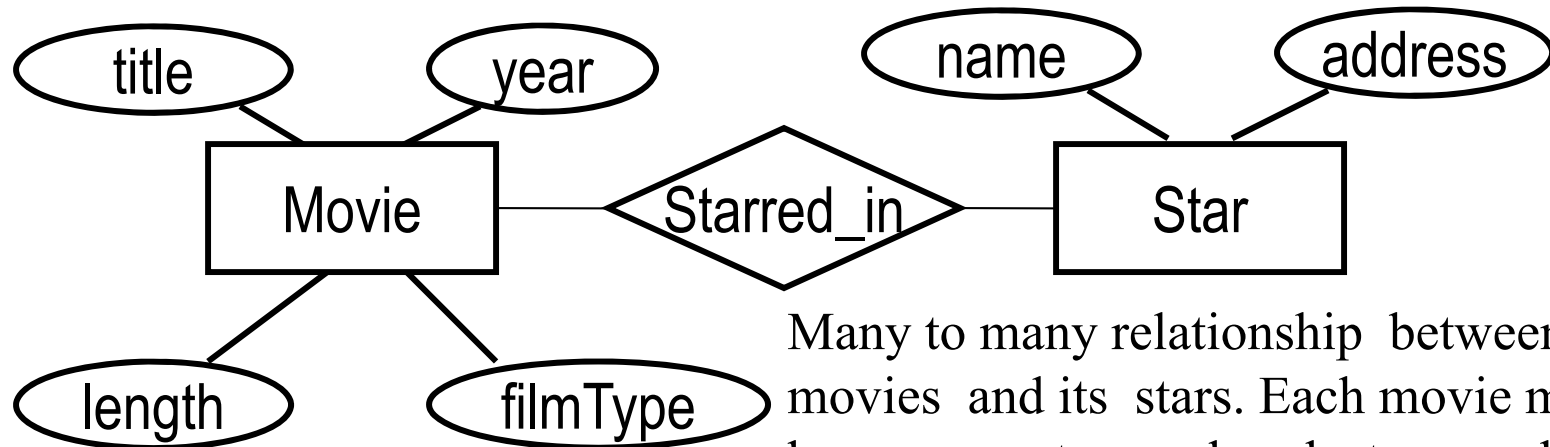
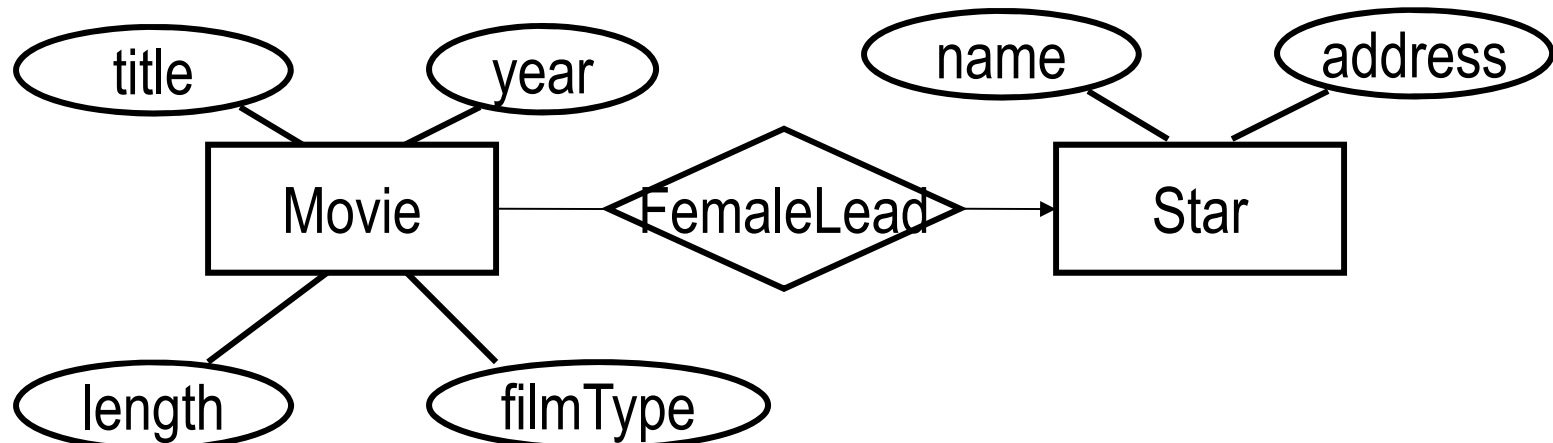
Alternate way of showing a many-to-one relationship

## Entities and entity sets



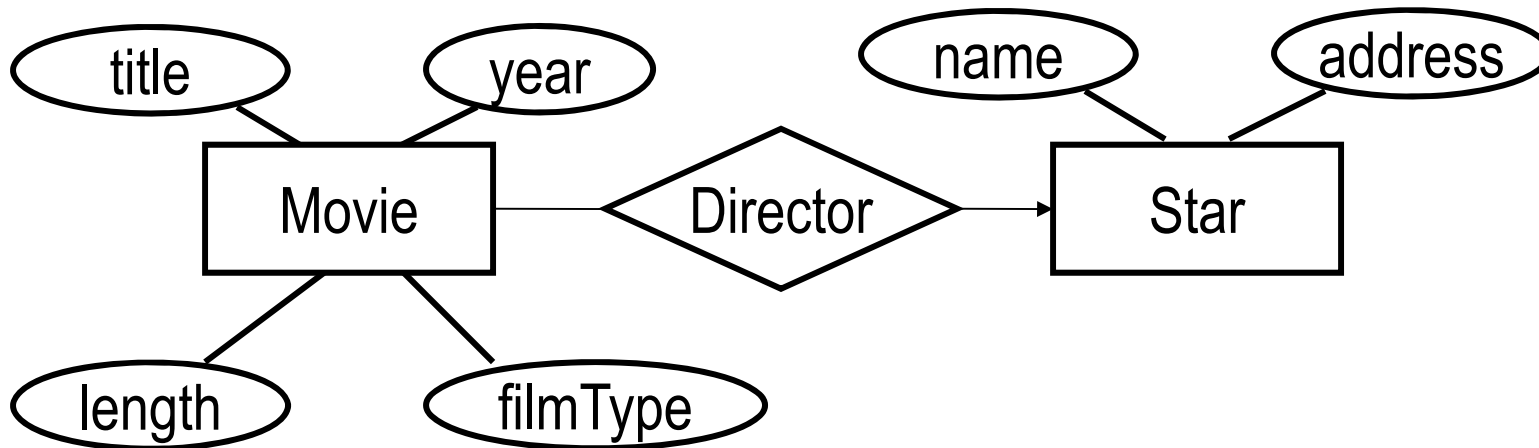
One to many relationship between a movie and its male lead (“hero”): Indicated by a arrow pointing to the “one side” – **A movie has but one main role, The star may be a lead in many movies**

## Entities and entity sets



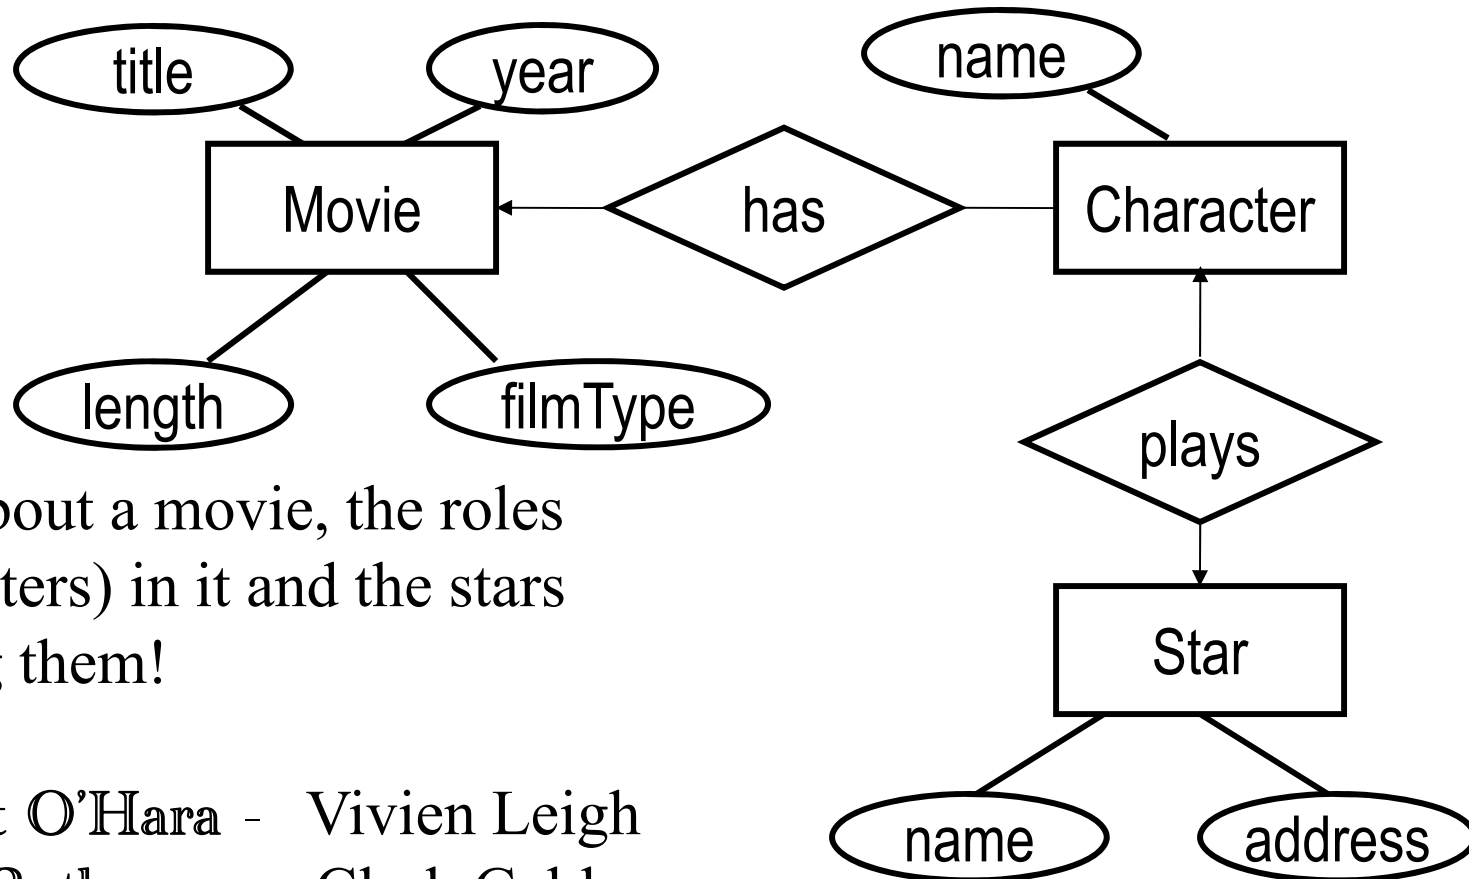
Many to many relationship between movies and its stars. Each movie may have many stars and each star may have featured in many movies. Indicated by no arrows on the connecting lines.

## Entities and entity sets



How about a movie and the roles(characters) in it and the stars playing them!

## Entities and entity sets



How about a movie, the roles (characters) in it and the stars playing them!

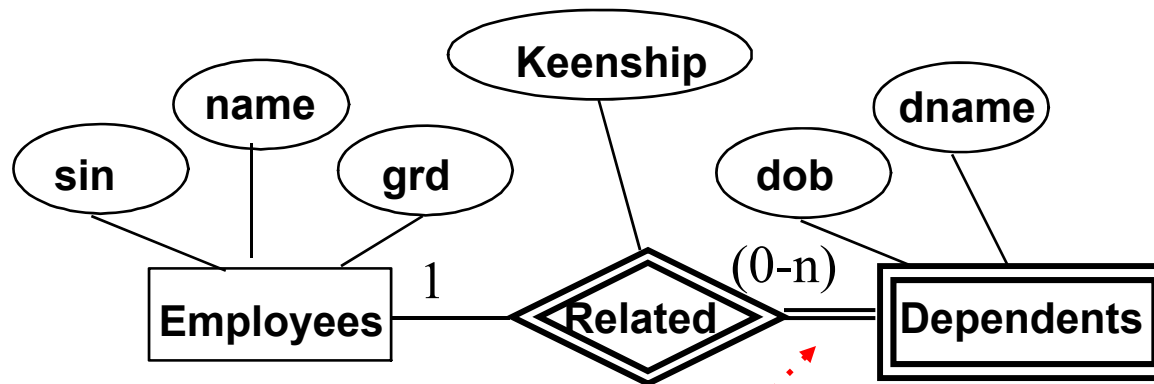
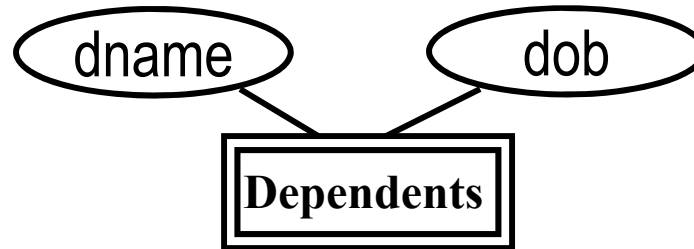
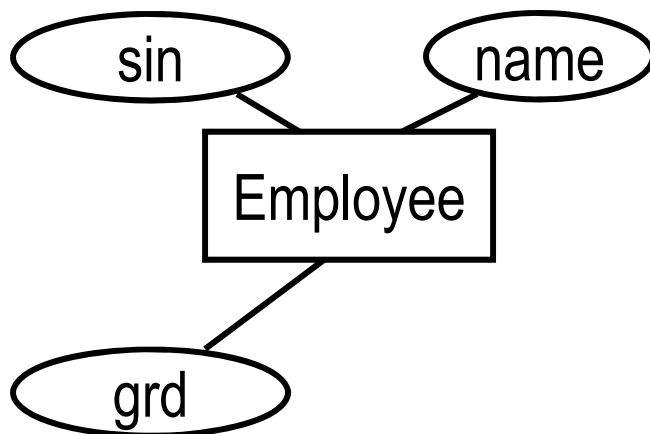
Scarlett O'Hara - Vivien Leigh

Rhett Butler - - Clark Gable

Alex Guinness plays eight members of the D'Ascoyne family in Kind hearts and coronets(1949)

Matt Damon played the lead in the *Bourne* trilogy.

# Entities and entity sets



An employee may have 0 to n dependents

*Total participation*

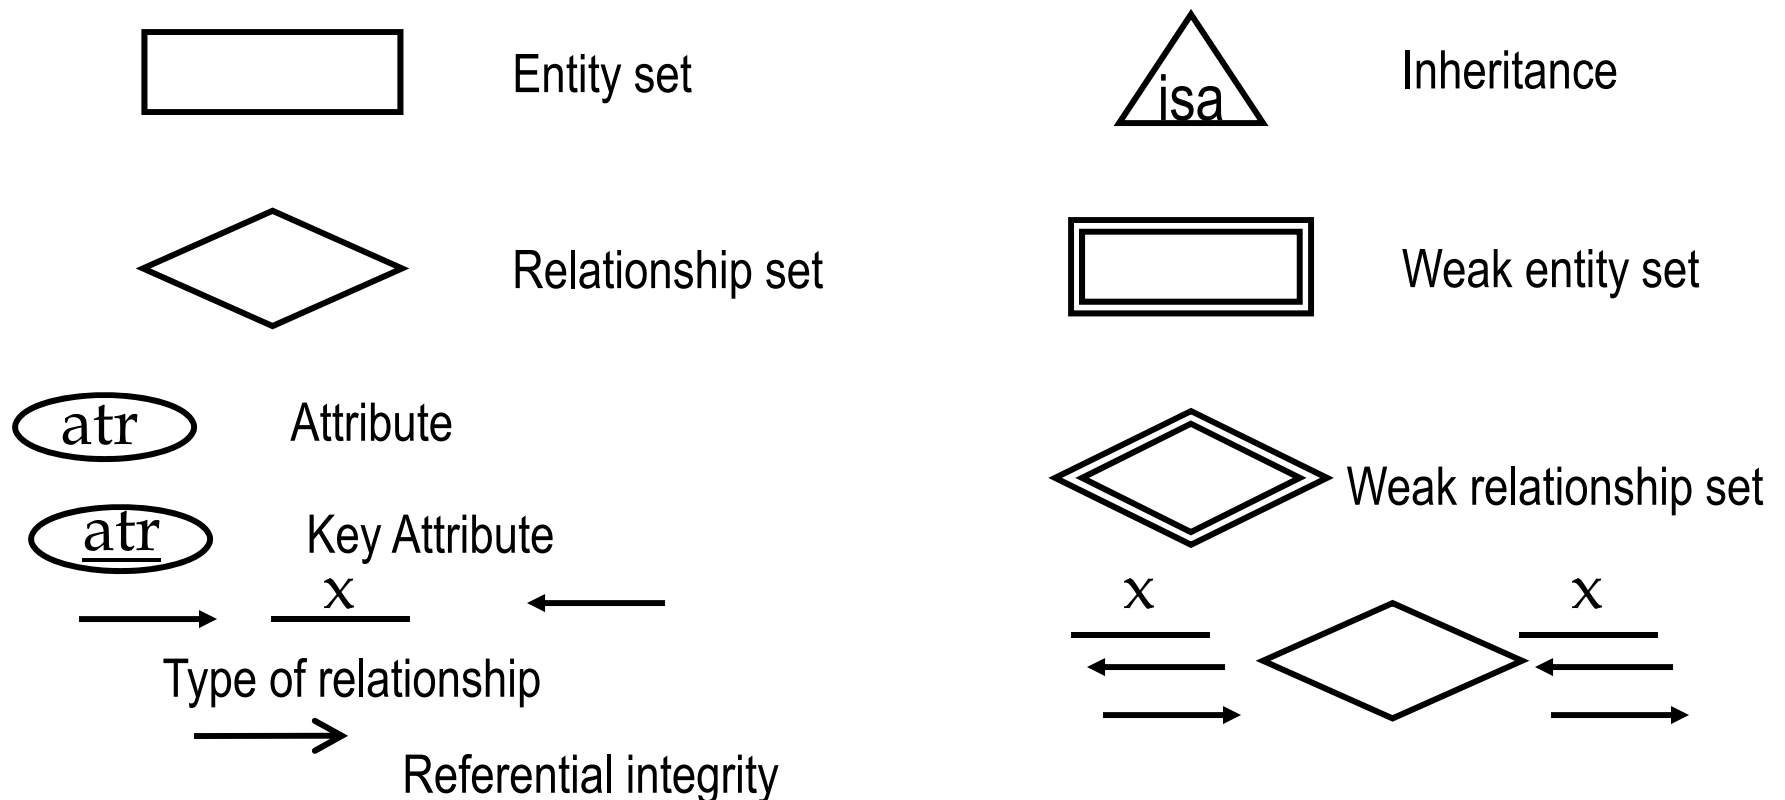
All dependents must be related to some employee (but only one!)

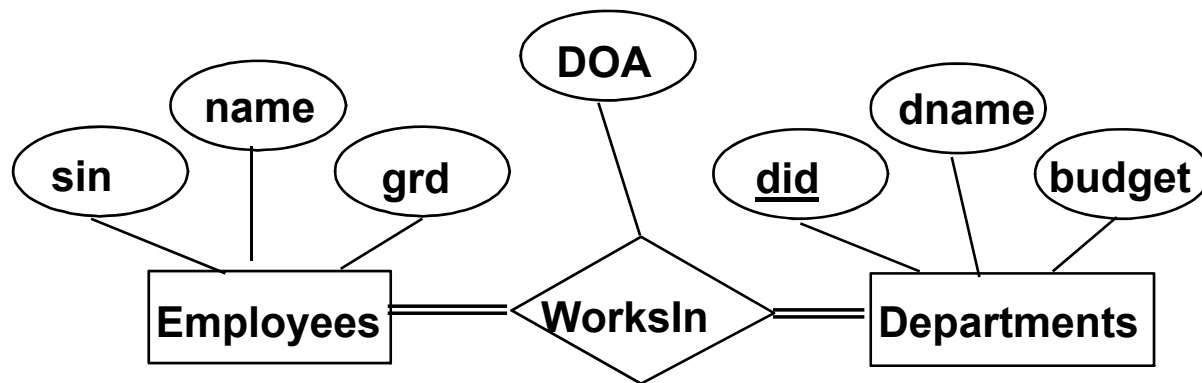
E/R model is a *graphical* approach to database modeling

E/R is widely used in database design

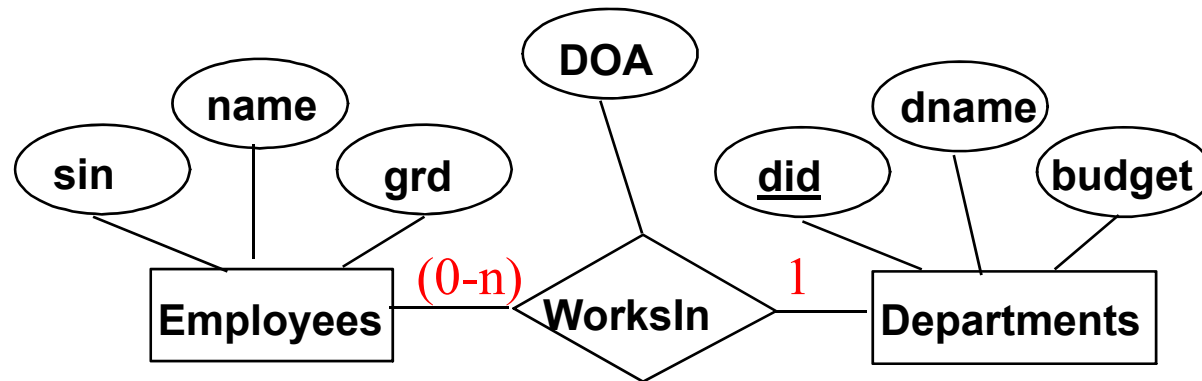
E/R model grew out of modeling application database

No standard for E/R diagrams: a number of variations

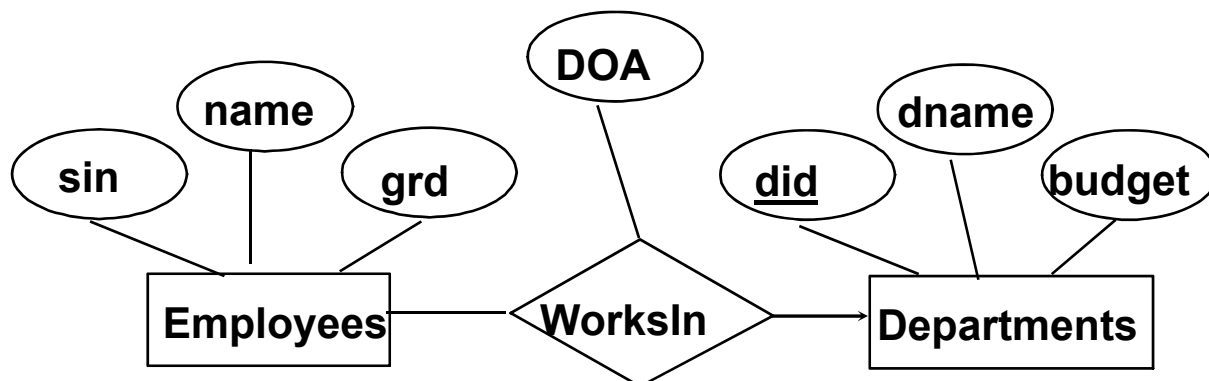




Total participation of all employees & departments  
In the WorksIn relationship



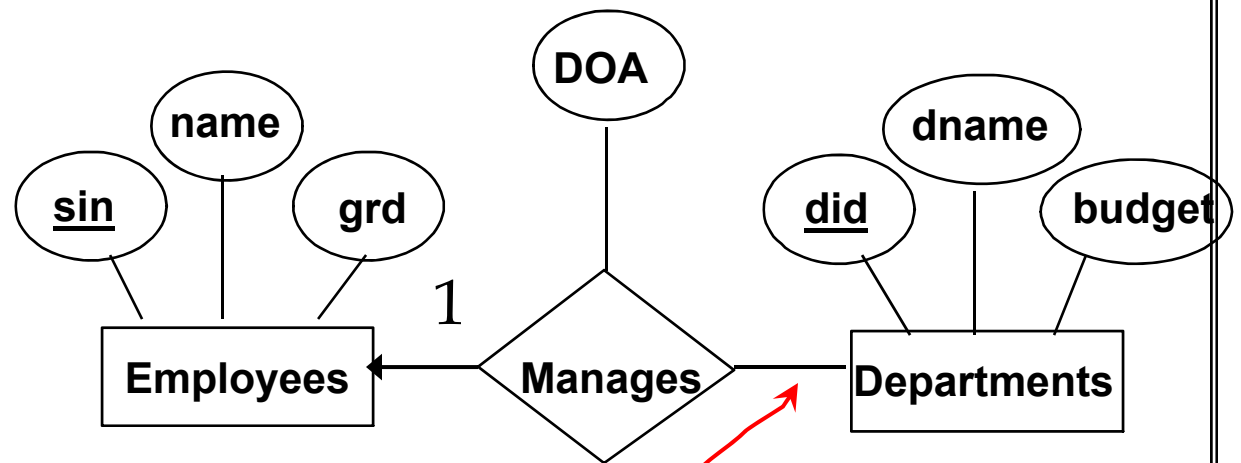
Many to one relationship:  
many employees in a  
department;  
but an employee is  
assigned to only one  
department



Alternate way of showing a many-to-one relationship



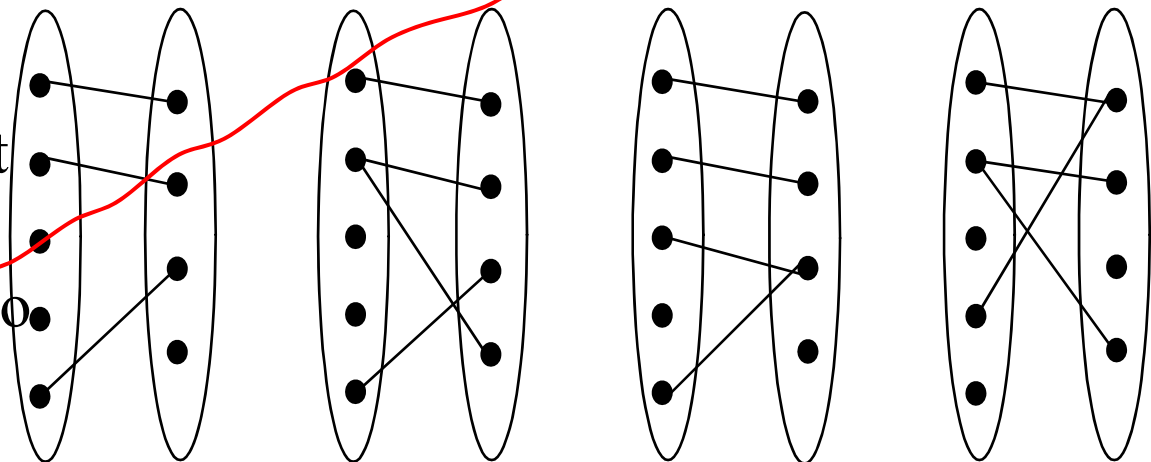
# Key Constraints



- ❖ Consider Works\_In:  
An employee can work in many departments; a dept can have many employees.

A department can have only one manager, an employee could manage many departments.

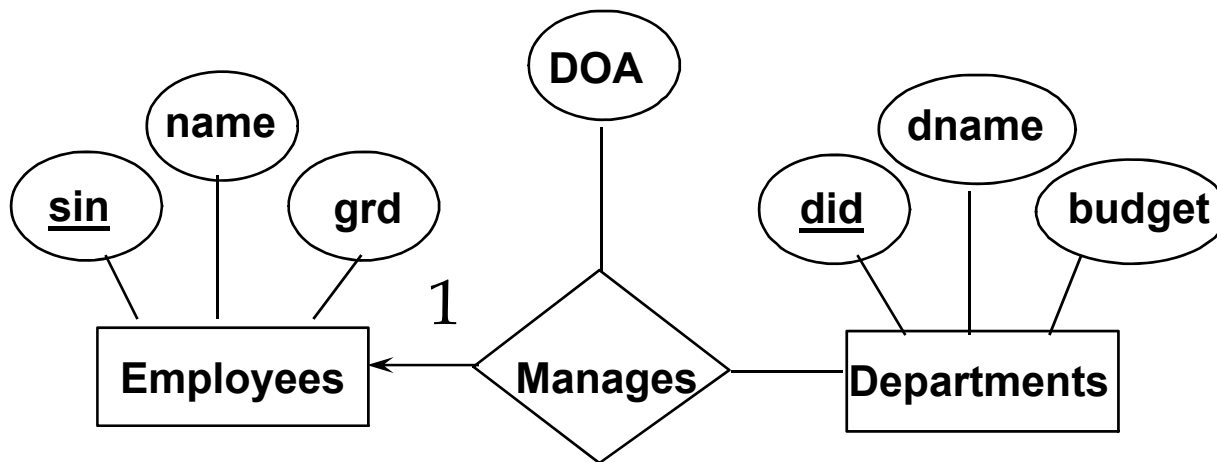
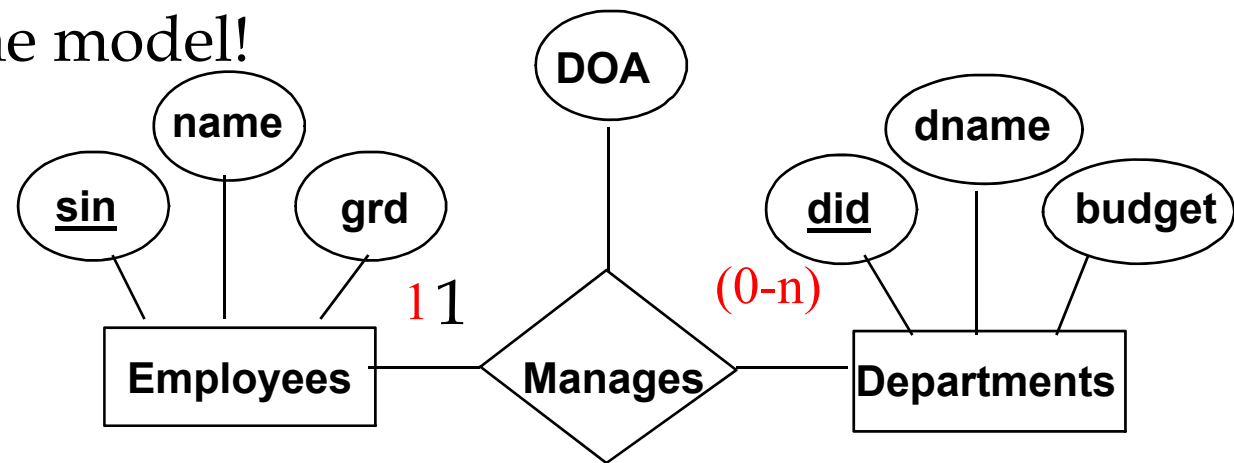
- ❖ In contrast, each dept has at most one manager, according to the key constraint on Manages.



- ❖ Relationship sets can have attributes
- ❖ In translating a relationship set to a relation, attributes of the relation must include:
  - ◆ Keys for each participating entity set (as foreign keys).
    - ✧ This set of attributes forms **superkey** for the relation.
  - ◆ All descriptive attributes.

```
CREATE TABLE WorksIn
(sin CHAR(9),
 did INTEGER,
 DOA DATE,
 PRIMARY KEY (sin, did),
 FOREIGN KEY (sin)
 REFERENCES Employees,
 FOREIGN KEY (did)
 REFERENCES Departments)
```

Alternate methods of  
showing the same model!



❖ Map relationship to a table:

◆ Note that **did** is the key now!

◆ Separate tables for Employees and Departments.

❖ Since each department has a unique manager, we could instead combine Manages and Departments.


CREATE TABLE Manages

( **sin** CHAR(9),  
**did** INTEGER,  
**DOA** DATE,  
**PRIMARY KEY** (did),  
FOREIGN KEY (sin) REFERENCES Employees,  
FOREIGN KEY (did) REFERENCES Departments)

CREATE TABLE DeptMgr

(**did** INTEGER, **sin** CHAR(9),  
dname CHAR(20),  
budget REAL,  
**DOA** DATE,  
**PRIMARY KEY** (did),  
**FOREIGN KEY** (sin) REFERENCES Employees)

Null for Dept. w/o manager!



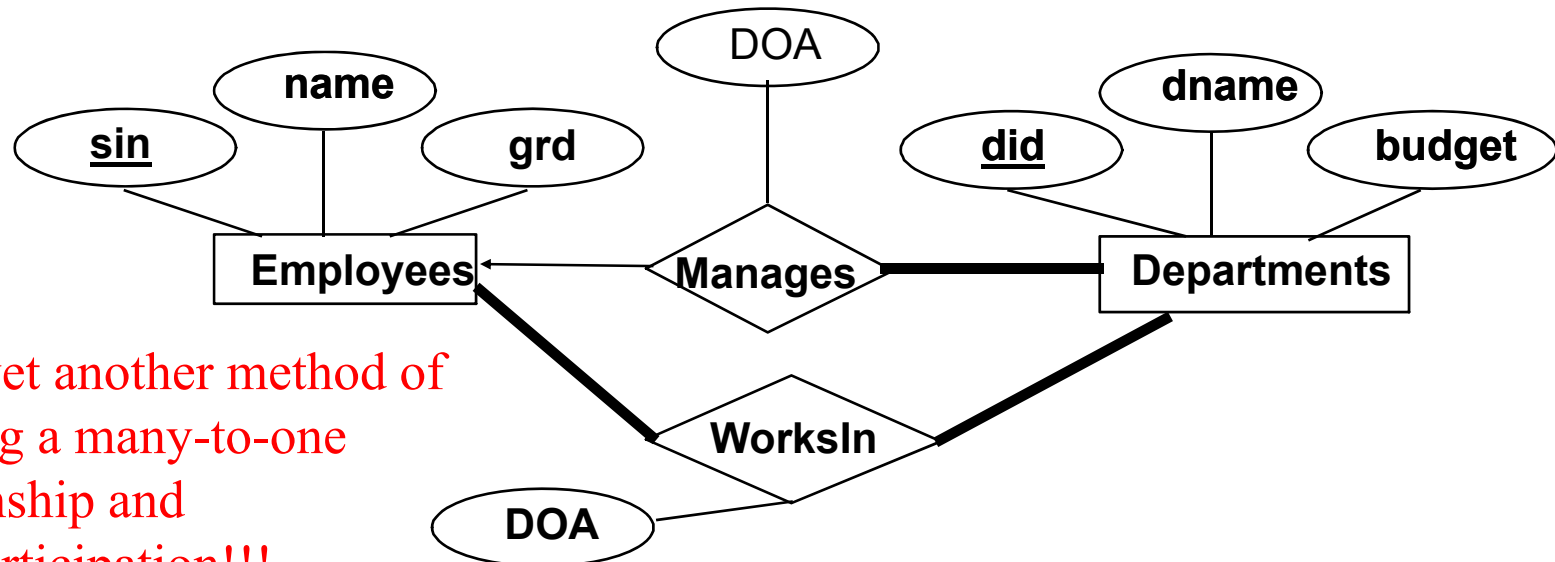
# Participation Constraints

❖ Every department has a manager (a business rule)  $\Rightarrow$

*participation constraint*:

◆ The participation of Departments in Manages is *total* (vs. *partial* i.e., *not all employees are managers*).

✧ Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *sin* value!)

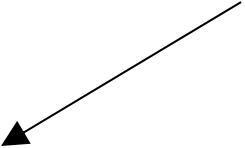


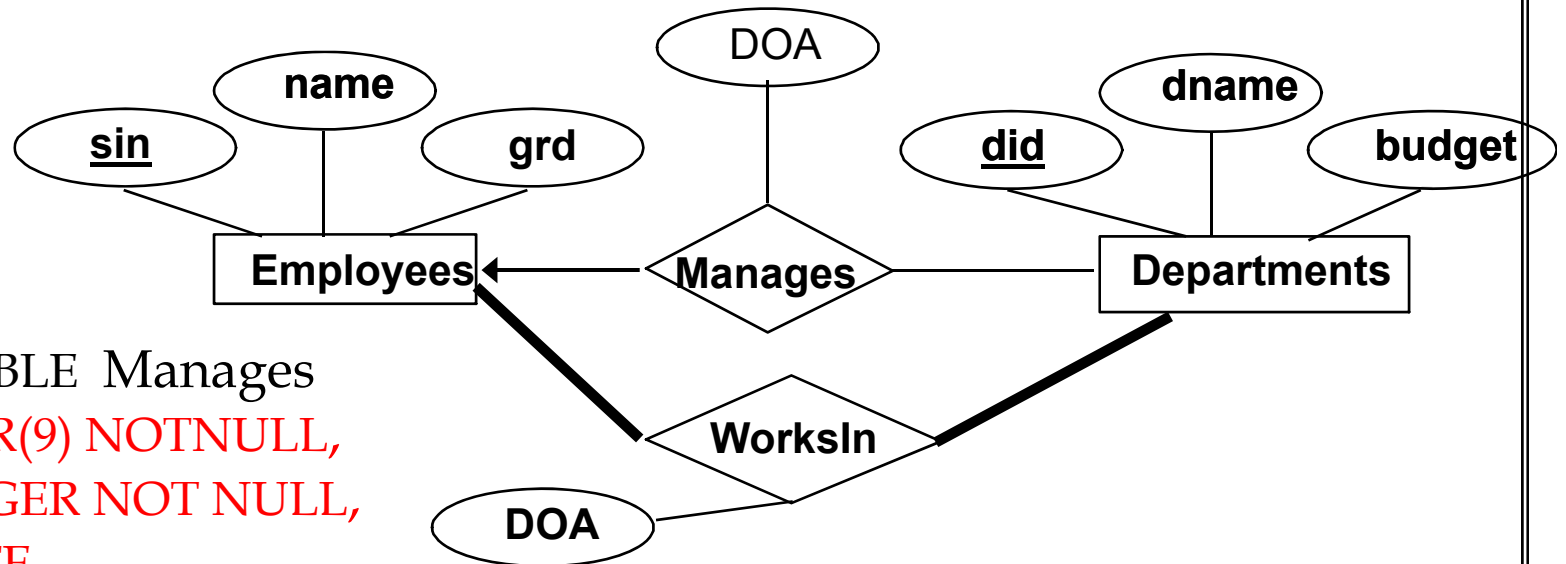
Note: yet another method of showing a many-to-one relationship and total participation!!!

## Participation Constraints: SQL

- ❖ A participation constraints involving one entity set in a binary relationship, can be expressed as follows without resorting to CHECK constraints.

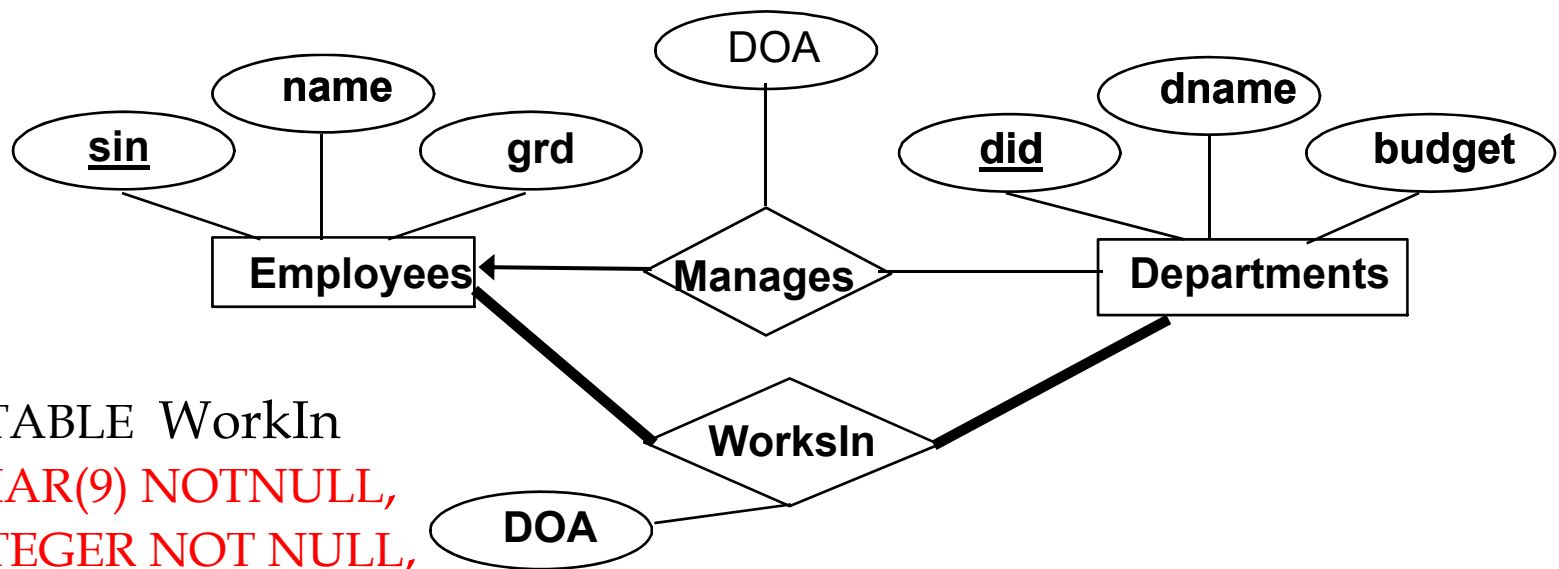
```
CREATE TABLE DeptMgr    Every department must
  (did INTEGER,          have a manager!
   dname CHAR(20),
   budget REAL,
   sin CHAR(9) NOT NULL,
   DOA DATE,
   PRIMARY KEY (did),
   FOREIGN KEY (sin) REFERENCES Employees,
   ON DELETE NO ACTION)
```





```
CREATE TABLE Manages
( sin CHAR(9) NOTNULL,
  did INTEGER NOT NULL,
  DOA DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (sin) REFERENCES Employees,
  FOREIGN KEY (did) REFERENCES Departments)
```

Here a department can exist without a manager.  
 We couldn't insert an occurrence of this relation  
 without having an occurrence of a department  
 and employee! Once inserted, does firing the manager  
 create problems?

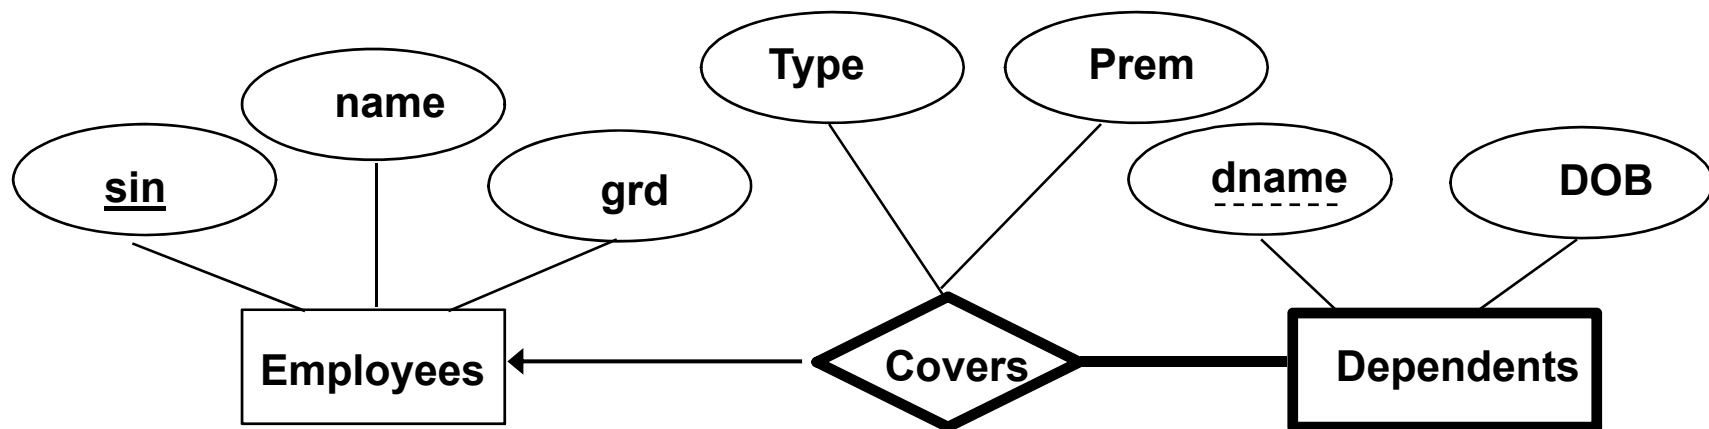


```
CREATE TABLE WorkIn
( sin CHAR(9) NOTNULL,
  did INTEGER NOT NULL,
  DOA DATE)
```

To ensure total participation of department in WorkIn, each did value must be in at least one tuple of WorkIn:  
enforced by assertion



- ❖ A *weak entity* can be identified uniquely only by considering the primary key of another *strong-owner* entity.
  - ◆ Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
  - ◆ Weak entity set must have total participation in this *identifying* relationship set.



- ❖ Weak entity set and identifying relationship set are translated into a single relation.
  - ◆ Weak entity  $\Rightarrow$  total participation
  - ◆ When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Covers
( dname CHAR(20),
  DOB DATE,
  Type INTEGER,
  Cost FLOAT,
  sin CHAR(9) NOT NULL,
  PRIMARY KEY (dname, sin),
  FOREIGN KEY (sin) REFERENCES Employees,
  ON DELETE CASCADE)
```

## Generalization, Specialization

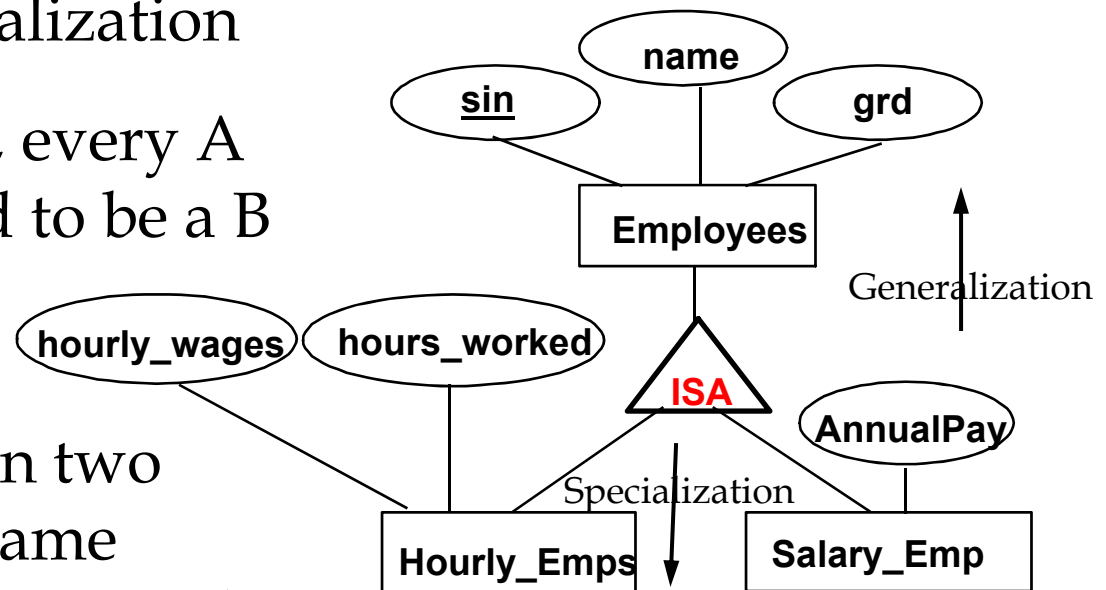
❖ If we declare A **ISA** B, every A entity is also considered to be a B entity. However, not implemented always:

❖ *Overlap constraints*: Can two subclasses contain the same instance of an entity? Can Carole be an Hourly\_Emps as well as a Salary\_Emps?

*(Allowed/disallowed)*

❖ Reasons for using ISA relationship:

- ◆ To add attributes specific to a subclass.
- ◆ To identify subset of an entity set that participate in a relationship.

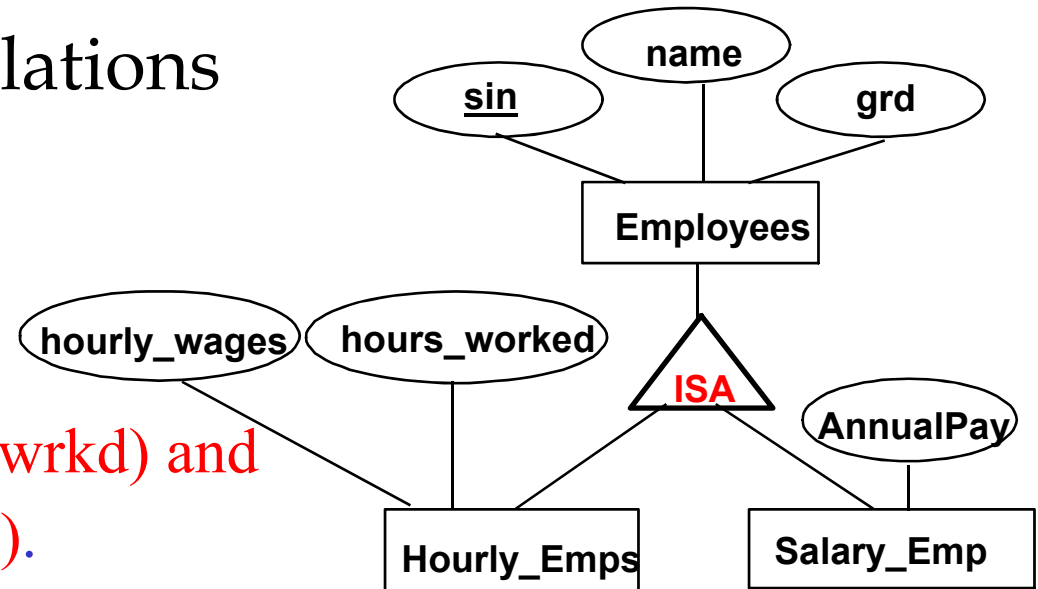


*Covering constraints*: Does every Employees entity also have to be an Hourly\_Emps or a Contract\_Emps entity? *(Yes/no)*

# ISA relationship to Relations

Create 3 relations:

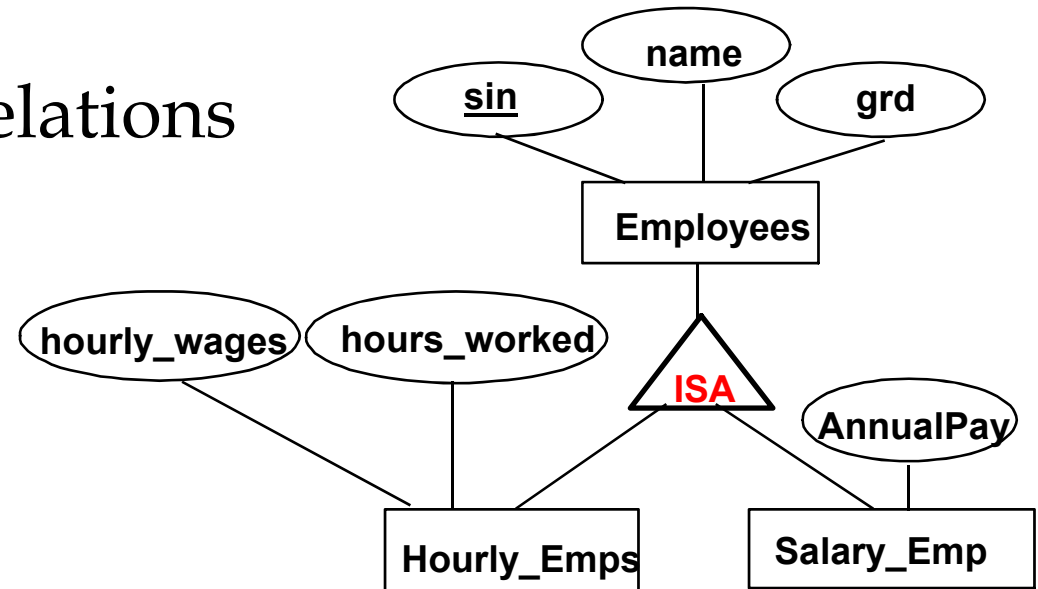
Employees(Sin, Name, Grd),  
Hourly\_Emps(Sin, Hwage, Hwrkd) and  
Salary\_Emps(Sin, AnnualPay).



❖ Every employee is recorded in Employees. For hourly employees, value for additional attribute are recorded in Hourly\_Emps; if referenced Employees tuple is deleted, Hourly\_Emps tuple must also be deleted.

◆ Queries involving all employees easy, those involving just Hourly\_Emps require a join with Employee to get inherited attributes.

# ISA relationship to Relations



❖ Create two relations:

❖ Hourly\_Emps(Sin, Name, Grd, HWrkd, Hwages) and Salary\_Emps (Sin, Name, Grd, AnnualPay).

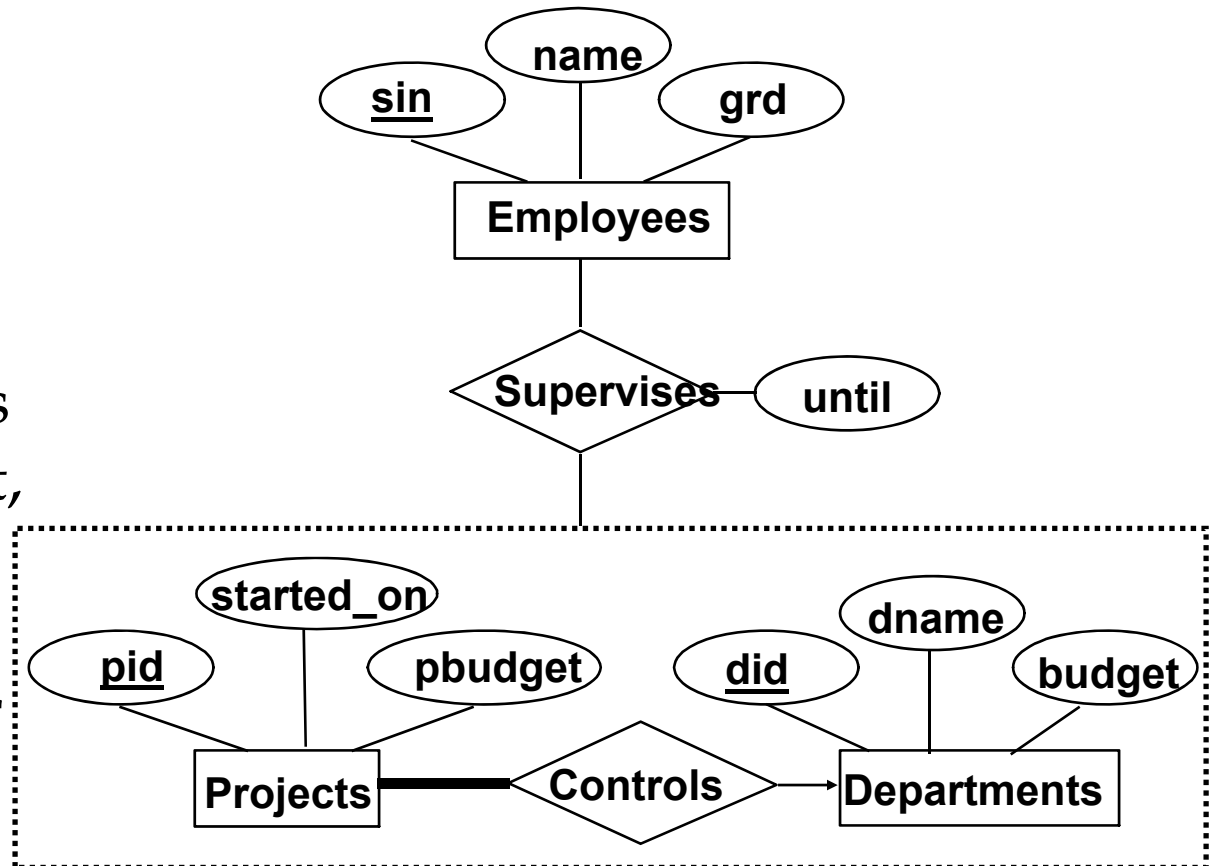
Each employee must be in one of these two subclasses.

*All employees require accessing Two relations*

# Aggregation

❖ Aggregation: models a relationship, involving entity sets and a *relationship* set, as an entity set. The aggregated entity participates in other relationships.

◆ Supervise mapped to table like any other relationship set.

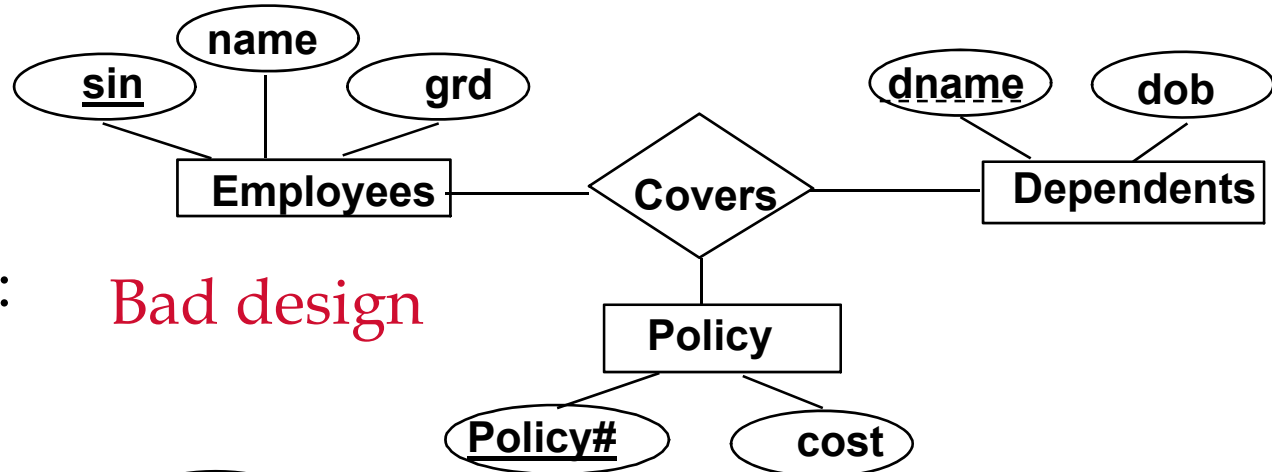


✗ Aggregation vs. ternary relationship:  
❖ Supervises is a distinct relationship, with its attribute.

## Binary vs. Ternary Relationships

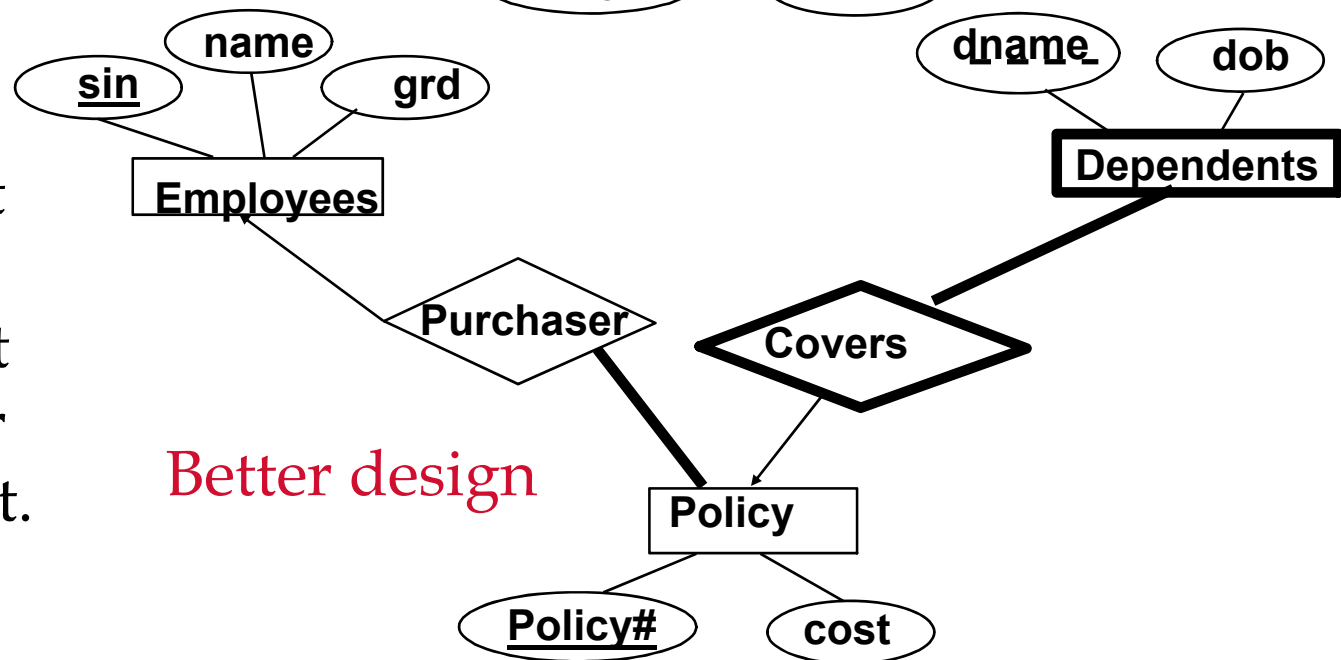
- ❖ If each policy is owned by just ONE employee:

Bad design



- ❖ Key constraint on Policy requires that it can only cover one dependent.

Better design



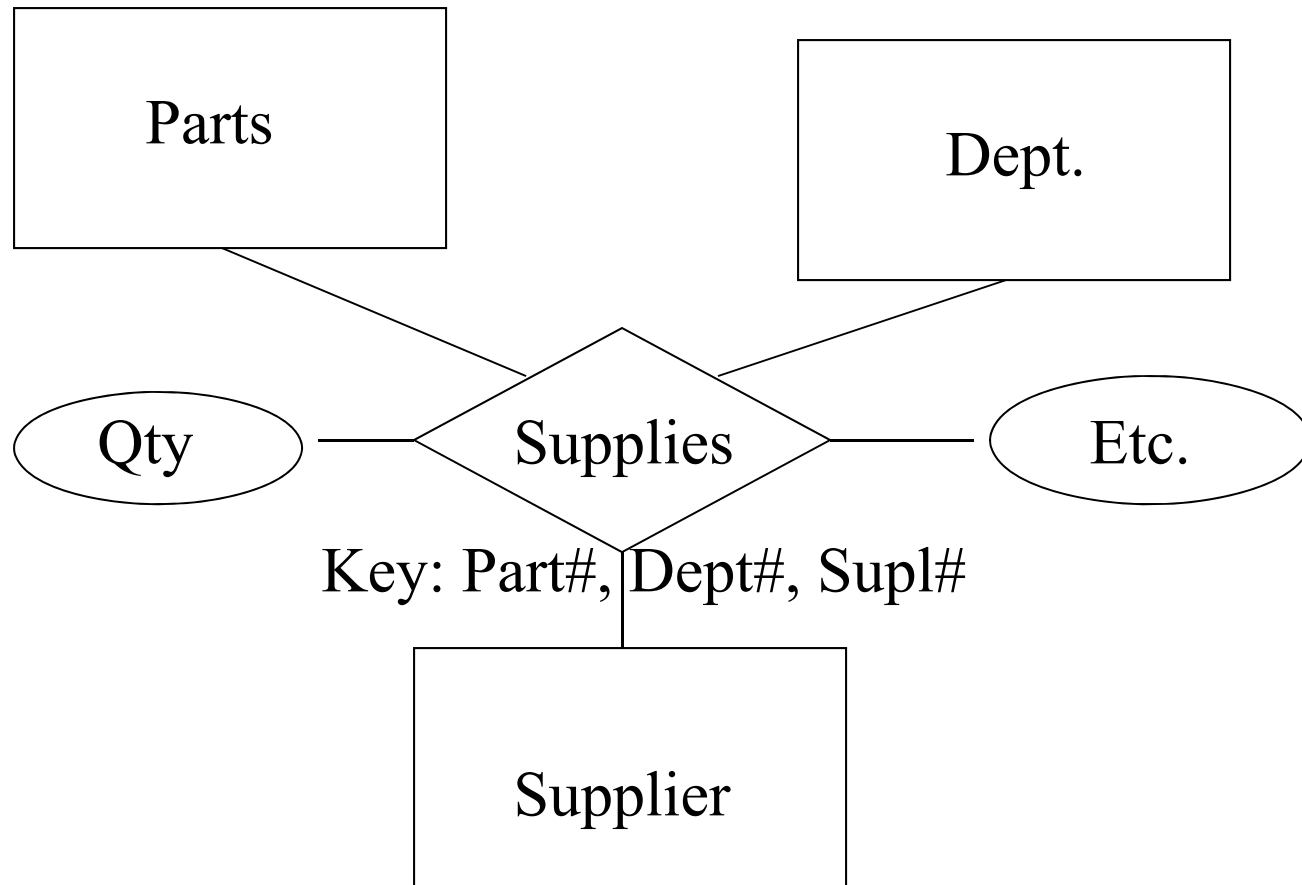
- ❖ The key constraints allow us to combine Purchaser with Policy and Covers with Dependents.
- ❖ Participation constraints lead to **NOT NULL** constraints.

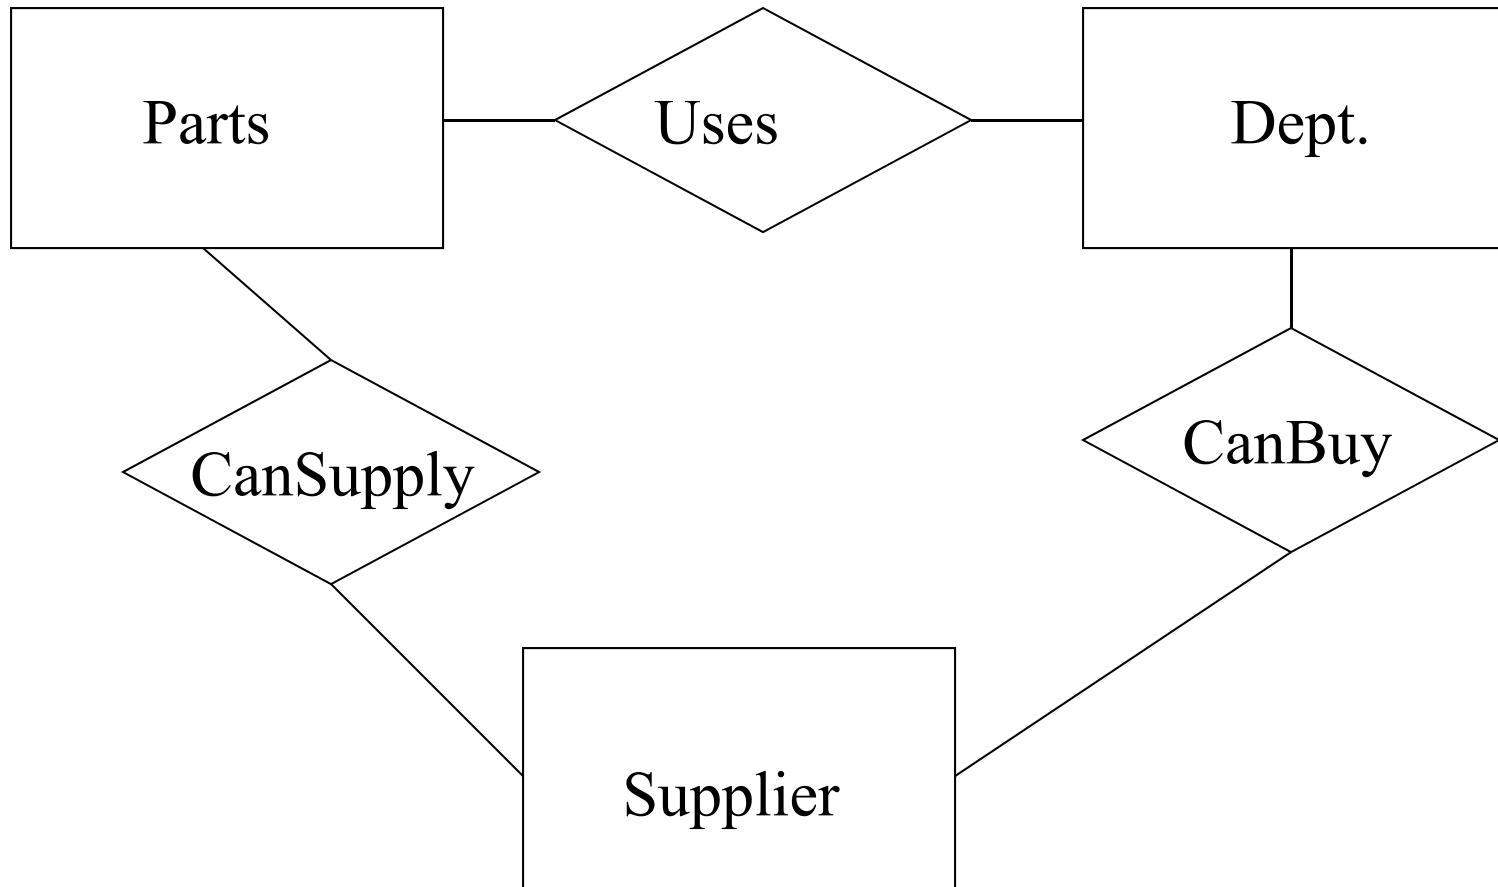
```
CREATE TABLE Policy (  
  policy# INTEGER,  
  cost REAL,  
  sin CHAR(9) NOT NULL,  
  PRIMARY KEY (policy#).  
  FOREIGN KEY (sin) REFERENCES  
    Employees,  
  ON DELETE CASCADE)
```

```
CREATE TABLE Dependents (  
  dname CHAR(20),  
  dob DATE,  
  policy# INTEGER,  
  PRIMARY KEY (dname, policy#).  
  FOREIGN KEY (policy#)  
    REFERENCES Policy,  
  ON DELETE CASCADE)
```



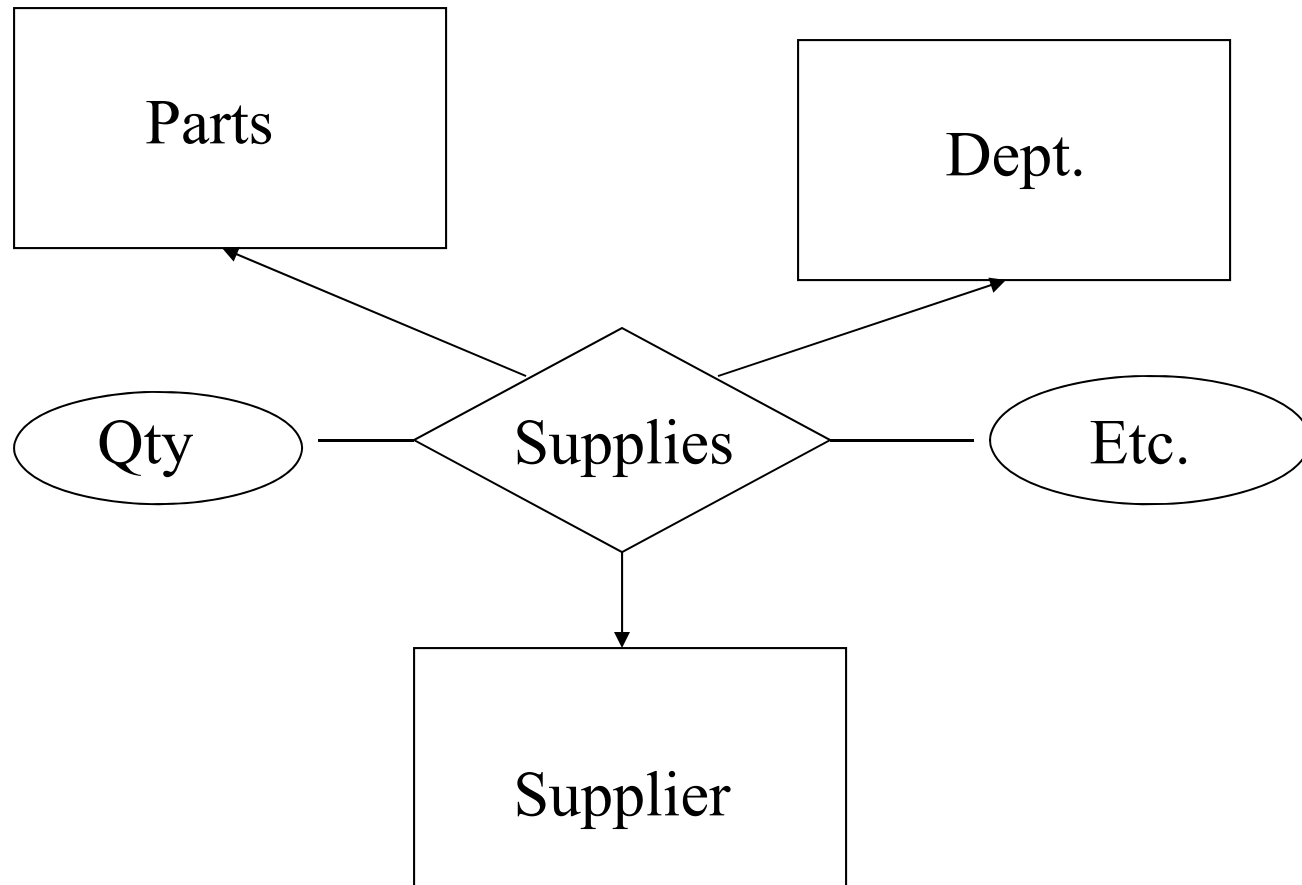
An example where a ternary relation is better than a number of binary relations is the following:



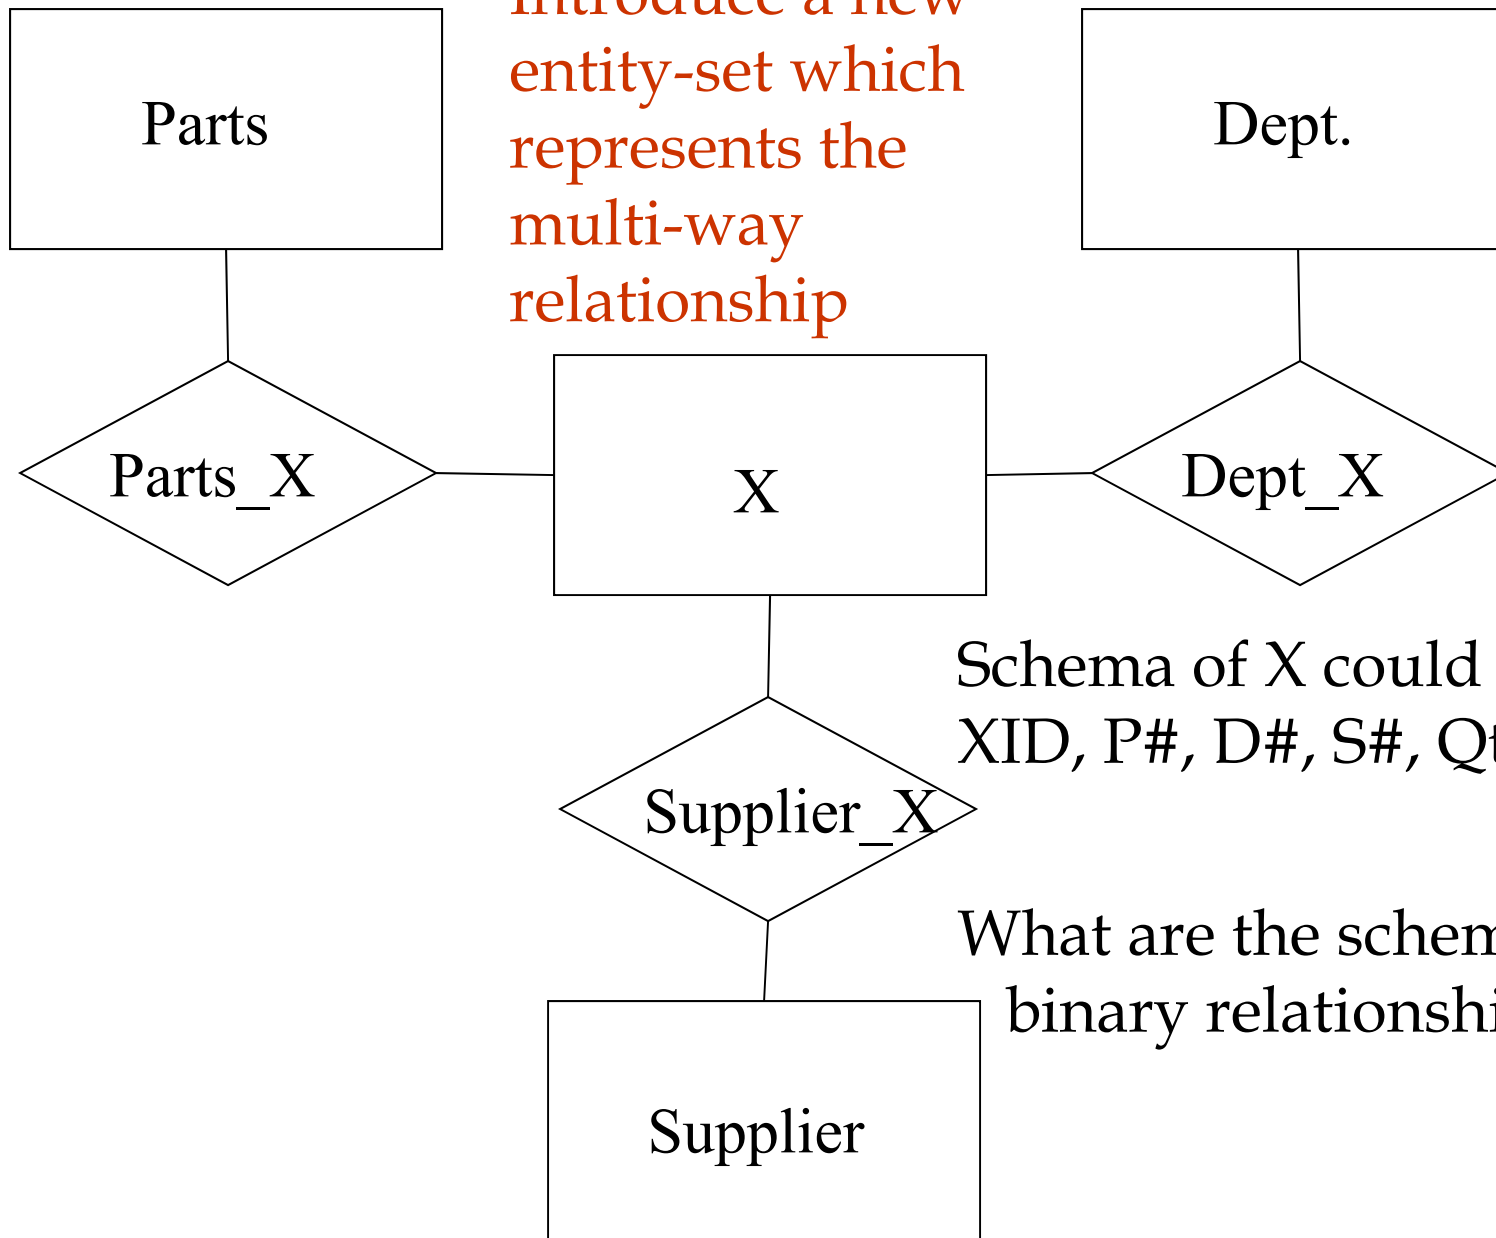


- ❖ **Supplies** relates entity sets **Parts**, **Departments** and **Suppliers**, and has descriptive attributes price, quantity etc.
- ❖ A number of binary relationships may not convey the semantics
- ❖ Supplier ``CanSupply`` Part, Dept. ``Uses`` Part, and Dept. ``Can Buy`` from Supplier does not imply that Dept has a PO to buy Part from S.
  - ◆ How do we record the following: which part, quantity price?

A department can order only one part from a supplier?



Introduce a new  
entity-set which  
represents the  
multi-way  
relationship



Schema of X could be  
XID, P#, D#, S#, Qty, ...

What are the schema of the  
binary relationship?

## Constraints Beyond the ER Model

### ❖ Functional dependencies:

- ◆ *A department can order only one part from a given supplier.*
  - ✧ Can't express this in ternary Supplies relationship.
- ◆ Normalization refines ER design by considering FDs.

### ❖ Inclusion dependencies:

- ◆ Special case: Foreign keys (ER model can express these).
- ◆ *e.g., At least 1 person must report to each manager. (Set of *sin* values in Manages must be subset of *supervisor\_ssn* values in Reports\_To.)*

### ❖ General constraints:

- ◆ *e.g., Manager's discretionary budget less than 10% of the combined budget of all departments he or she manages.*